

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum- 590014, Karnataka.



## LAB REPORT on **Machine Learning (23CS6PCMAL)**

*Submitted by*

**Roshni P(1BM22CS223)**

*in partial fulfillment for the award of the degree of*

## **BACHELOR OF ENGINEERING**

*in*

## **COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING**  
(Autonomous Institution under VTU)  
**BENGALURU - 560019**  
**February 2025 – July 25**

# B.M.S. College of Engineering

Bull Temple Road, Bangalore 560019

(Affiliated To Visvesvaraya Technological University, Belgaum)

## Department of Computer Science and Engineering



### CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Roshni P(1BM22CS223)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfilment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Laboratory report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Saritha A N  
Assistant Professor  
Department of CSE, BMSCE

Dr. Kavitha Sooda  
Professor & HOD  
Department of CSE, BMSCE

## INDEX

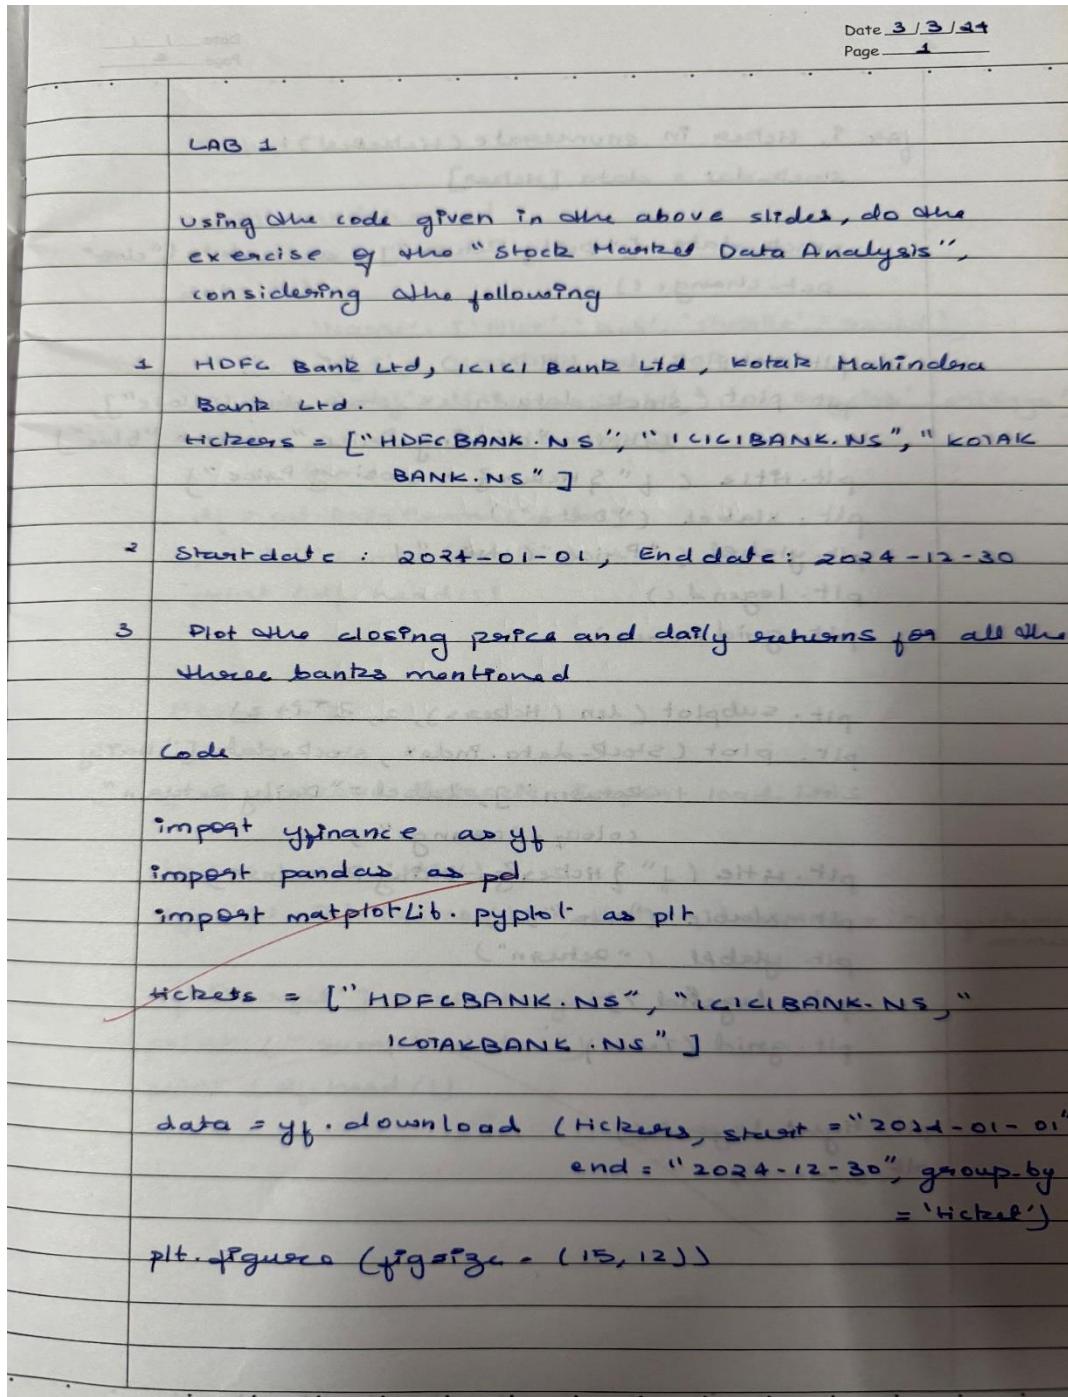
<b>Sl. No.</b>	<b>Date</b>	<b>Experiment Title</b>	<b>Page No.</b>
1	03.03.25	Write a python program to import and export data using pandas library functions.	1
2	10.03.25	Demonstrate various data pre-processing techniques for a given dataset.	5
3	24.03.25	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	10
4	17.03.25	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset.	14
5	24.03.25	Build Logistic Regression Model for a given dataset.	18
6	07.04.25	Build KNN Classification model for a given dataset.	22
7	21.04.25	Build Support vector machine model for a given dataset.	27
8	05.05.25	Implement Random forest ensemble method on a given dataset.	32
9	05.05.25	Implement Boosting ensemble method on a given dataset.	35
10	05.05.25	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	40
11	05.05.25	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	45

Github Link: <https://github.com/RoshniP223/ML-LAB>

## LABORATORY PROGRAM – 1

Write a python program to import and export data using Pandas library functions

### OBSERVATION



Date / /  
Page 2

```

for i, tickers in enumerate(tickers):
    stock_dat = data[tickers]
    stock_data["Daily Return"] = stock_data["close"] - pct_change()

plt.subplot(len(tickers), 2, 2 * i + 1)
plt.plot(stock_data.index, stock_data["close"], label="closing Price", color="blue")
plt.title(f" {tickers} - closing Price")
plt.xlabel("Date")
plt.ylabel("Price (INR)")
plt.legend()
plt.grid(True)

plt.subplot(len(tickers), 2, 2 * i + 2)
plt.plot(stock_data.index, stock_data["Daily Return"], label="Daily Return", color="orange")
plt.title(f" {tickers} - Daily Returns")
plt.xlabel("Date")
plt.ylabel("Return")
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

```

Date / /  
Page 3

**Method - 1**

```

import pandas as pd
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [25, 30, 35, 40],
    'city': ['New York', 'Los Angeles', 'Chicago',
             'Houston']
}
df = pd.DataFrame(data)
print("sample data:")
print(df.head())

```

**Method - 2**

```

from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("sample data:")
print(df.head())

```

### Method 3

```
file-path = 'data.csv'  
df = pd.read_csv(file-path)  
print("Sample data :")  
print(df.head())  
print("\n")
```

### Method 4

```
df = pd.read_csv('mobiles-dataset-2025.csv')  
print("Sample data :")  
print(df.head())
```

### Output

sample data :

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago
3	David	40	Houston

Sample data

	sepal length	sepal width	petal length	petal width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
0	0			
1	0			

## CODE WITH OUTPUT

### Diabetes Dataset

```
df=pd.read_csv('/content/Dataset of Diabetes .csv')
df.head()
```

	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS	
0	502	17975	F	50	4.7	46		4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
1	735	34221	M	26	4.5	62		4.9	3.7	1.4	1.1	2.1	0.6	23.0	N
2	420	47975	F	50	4.7	46		4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
3	680	87656	F	50	4.7	46		4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
4	504	34223	M	33	7.1	46		4.9	4.9	1.0	0.8	2.0	0.4	21.0	N

```
df.shape
```

```
(1000, 14)
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   ID          1000 non-null    int64  
 1   No_Pation   1000 non-null    int64  
 2   Gender       1000 non-null    object  
 3   AGE          1000 non-null    int64  
 4   Urea         1000 non-null    float64 
 5   Cr           1000 non-null    int64  
 6   HbA1c        1000 non-null    float64 
 7   Chol          1000 non-null    float64 
 8   TG            1000 non-null    float64 
 9   HDL           1000 non-null    float64 
 10  LDL           1000 non-null    float64 
 11  VLDL          1000 non-null    float64 
 12  BMI           1000 non-null    float64 
 13  CLASS         1000 non-null    object  
dtypes: float64(8), int64(4), object(2)
memory usage: 109.5+ KB
None
```

```

# Summary statistics
print(df.describe())

      ID   No_Pation      AGE      Urea      Cr \
count 1000.000000 1.000000e+03 1000.000000 1000.000000 1000.000000
mean  340.500000 2.705514e+05  53.528000  5.124743  68.943000
std   240.397673 3.380758e+06  8.799241  2.935165  59.984747
min   1.000000 1.230000e+02  20.000000  0.500000  6.000000
25%  125.750000 2.406375e+04  51.000000  3.700000  48.000000
50%  300.500000 3.439550e+04  55.000000  4.600000  60.000000
75%  550.250000 4.538425e+04  59.000000  5.700000  73.000000
max  800.000000 7.543566e+07  79.000000 38.900000 800.000000

      HbA1c      Chol      TG      HDL      LDL \
count 1000.000000 1000.000000 1000.000000 1000.000000 1000.000000
mean  8.281160  4.862820  2.349610  1.204750  2.609790
std   2.534003  1.301738  1.401176  0.660414  1.115102
min   0.900000  0.000000  0.300000  0.200000  0.300000
25%  6.500000  4.000000  1.500000  0.900000  1.800000
50%  8.000000  4.800000  2.000000  1.100000  2.500000
75% 10.200000  5.600000  2.900000  1.300000  3.300000
max  16.000000 10.300000 13.800000  9.900000  9.900000

      VLDL      BMI
count 1000.000000 1000.000000
mean  1.854700 29.578020
std   3.663599 4.962388
min   0.100000 19.000000
25%  0.700000 26.000000
50%  0.900000 30.000000
75%  1.500000 33.000000
max  35.000000 47.750000

missing_values=df.isnull().sum()
print(missing_values[missing_values > 0])

Series([], dtype: int64)

```

```

categorical_cols = df.select_dtypes(include=['object']).columns
print("Categorical columns identified:", categorical_cols)
if len(categorical_cols) > 0:
    df = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
    print("\nDataFrame after one-hot encoding:")
    print(df.head())
else:
    print("\nNo categorical columns found in the dataset.")

```

Categorical columns identified: Index(['Gender', 'CLASS'], dtype='object')

DataFrame after one-hot encoding:

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	\
0	502	17975	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	
1	735	34221	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	
2	420	47975	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	
3	680	87656	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	
4	504	34223	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	

	Gender_M	Gender_f	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y	
0	False	False	False	False	False	False	
1	True	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	True	False	False	False	False	False	

```

from sklearn.preprocessing import MinMaxScaler, StandardScaler
import pandas as pd

numerical_cols = df.select_dtypes(include=['number']).columns

scaler = MinMaxScaler()
df_minmax = df.copy() # Create a copy to avoid modifying the original
df_minmax[numerical_cols] = scaler.fit_transform(df[numerical_cols])

scaler = StandardScaler()
df_standard = df.copy()
df_standard[numerical_cols] = scaler.fit_transform(df[numerical_cols])
print("\nDataFrame after Min-Max Scaling:")
print(df_minmax.head())
print("\nDataFrame after Standardization:")
print(df_standard.head())

```

DataFrame after Min-Max Scaling:

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	\
0	0.627034	0.000237	0.508475	0.109375	0.050378	0.264901	0.407767	
1	0.918648	0.000452	0.101695	0.104167	0.070529	0.264901	0.359223	
2	0.524406	0.000634	0.508475	0.109375	0.050378	0.264901	0.407767	
3	0.849812	0.001160	0.508475	0.109375	0.050378	0.264901	0.407767	
4	0.629537	0.000452	0.220339	0.171875	0.050378	0.264901	0.475728	

	TG	HDL	LDL	VLDL	BMI	Gender_M	Gender_f	\
0	0.044444	0.226804	0.114583	0.011461	0.173913	False	False	
1	0.081481	0.092784	0.187500	0.014327	0.139130	True	False	
2	0.044444	0.226804	0.114583	0.011461	0.173913	False	False	
3	0.044444	0.226804	0.114583	0.011461	0.173913	False	False	
4	0.051852	0.061856	0.177083	0.008596	0.069565	True	False	

	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

DataFrame after Standardization:

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	\
0	0.672140	-0.074747	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
1	1.641852	-0.069940	-3.130017	-0.212954	-0.115804	-1.334983	-0.893730	
2	0.330868	-0.065869	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
3	1.412950	-0.054126	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
4	0.680463	-0.069939	-2.334096	0.673299	-0.382672	-1.334983	0.028576	

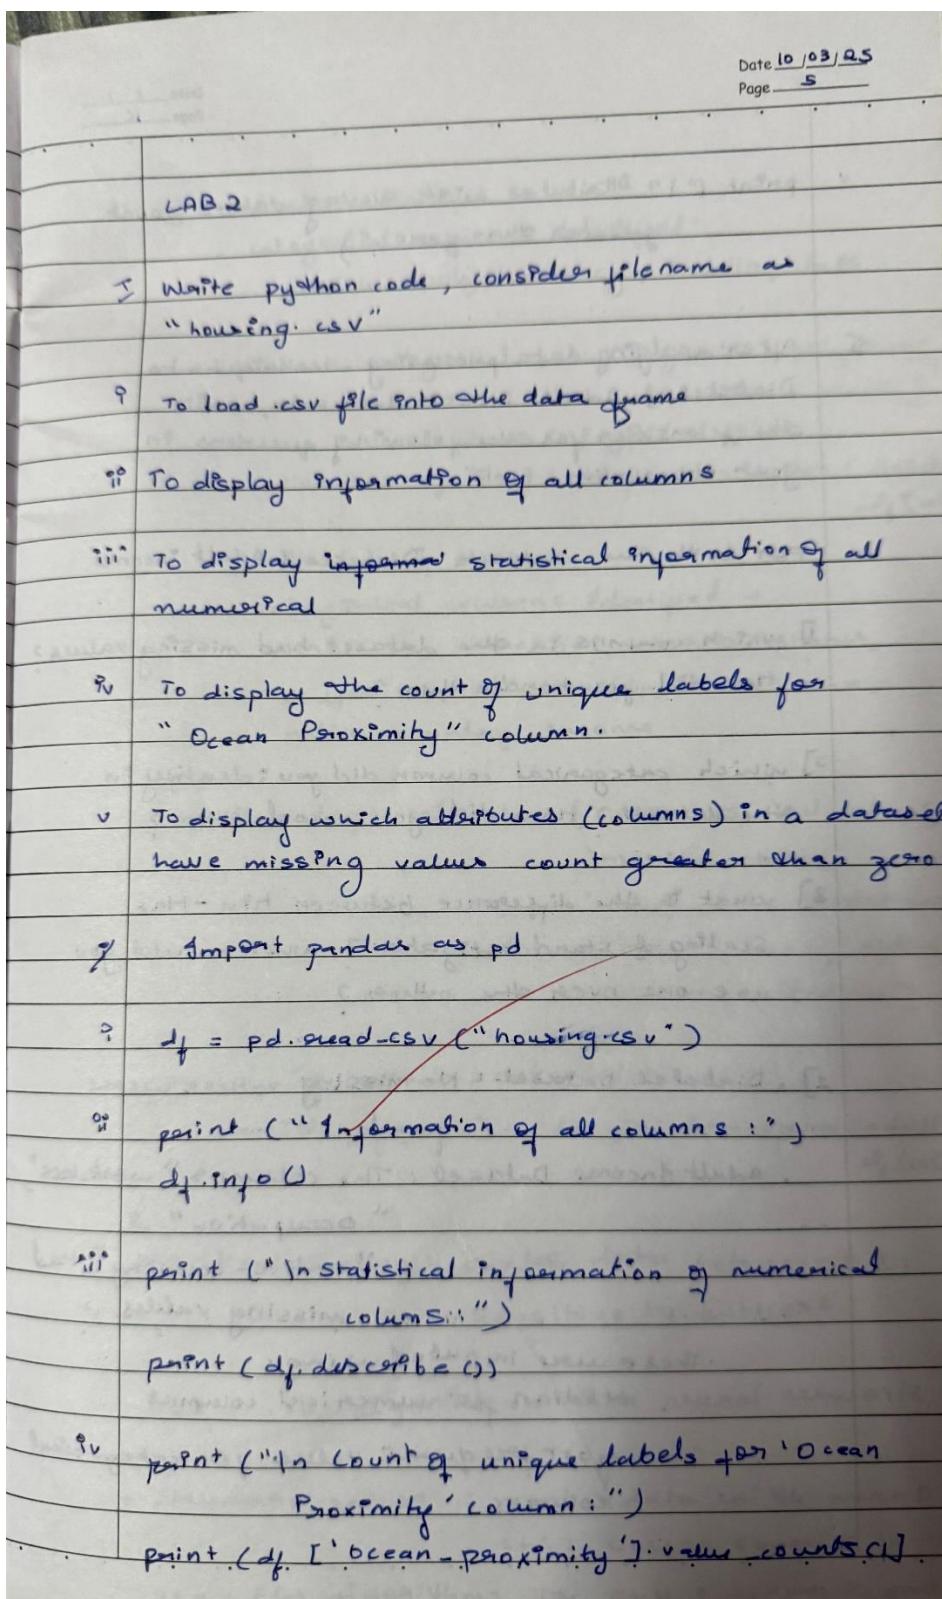
	TG	HDL	LDL	VLDL	BMI	Gender_M	Gender_f	\
0	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	False	False	
1	-0.678063	-0.158692	-0.457398	-0.342649	-1.326239	True	False	
2	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	False	False	
3	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	False	False	
4	-0.963680	-0.613180	-0.547121	-0.397267	-1.729472	True	False	

	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

## LABORATORY PROGRAM – 2

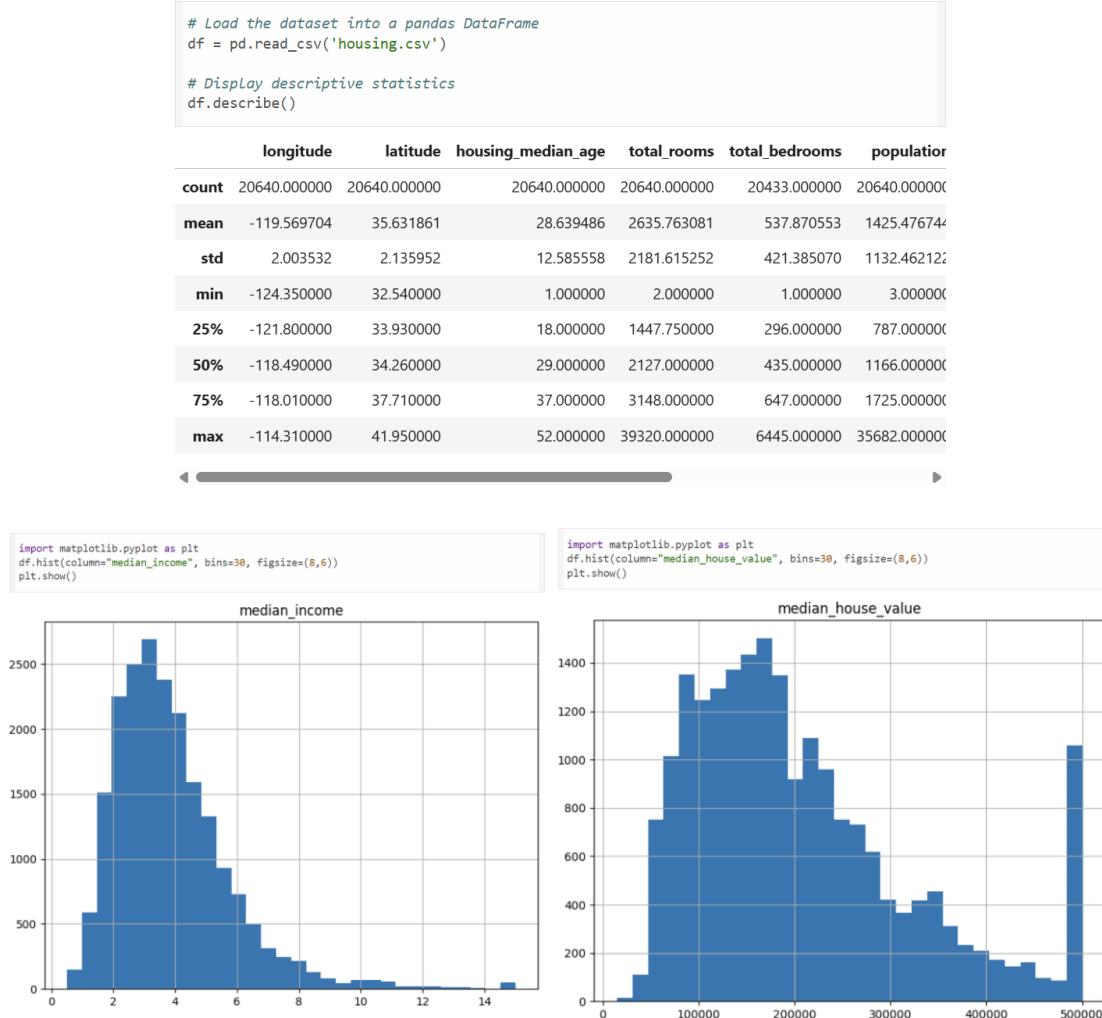
Demonstrate various data pre-processing techniques for a given dataset

### OBSERVATION BOOK



<p>v print ("{} Attributes with missing values count greater than zero.")</p> <p>ii After applying data processing techniques to Diabetes &amp; Adult income data sets, write the answers for the following questions in your observation book</p> <p>FDT both the datasets Diabetes &amp; Adult Income</p> <p>1] which columns in the dataset had missing values? How did you handle them?</p> <p>2] which categorical column did you identify in the dataset? How did you encode them?</p> <p>3] what is the difference between Min-Max Scaling &amp; Standardization? when would you use one over the other?</p> <p>1] Diabetes Dataset : No missing values were found.</p> <p>. Adult Income Dataset : The columns "workclass", "occupation", &amp; "native-country" had missing values.</p> <p>These were imputed using:</p> <ul style="list-style-type: none"> <li>Median for numerical columns</li> <li>most frequent value for categorical columns.</li> </ul>	<p>Date / / Page 6</p> <p>2] Diabetes Dataset :</p> <ul style="list-style-type: none"> <li>categorical columns identified : gender, polypharmacy, insulin, class</li> </ul> <p>diabetes - categorical - cols = ["gender", "polypharmacy", "insulin", "class"]</p> <p>for col in diabetes : categorical_cols.</p> <p>diabetes_df [col] = le.fit_transform (diabetes_df [col])</p> <p>Adult Income Dataset</p> <ul style="list-style-type: none"> <li>categorical columns identified - workclass, education, marital-status, occupation, relationship, race, sex, native-country, income</li> </ul> <p>adult - categorical - cols = ["workclass", "education", "marital-status", "occupation", "relationship", "race", "sex", "native-country", "income"]</p> <p>for col in adult - categorical - cols :</p> <p>adult_df [col] = le.fit_transform (adult_df [col])</p> <p>3] Min-Max Scaling : scales data between 0 &amp; 1 Sensitive to outliers</p> <p>Use : FDT algorithms like neural network/2</p> <p>Standardization : center data with mean 0 &amp; std 1. Less sensitive to outliers</p> <p>Use : FDT algorithms like SVM &amp; Linear Regression</p>
--	--

## CODE WITH OUTPUT



```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit

# Load the dataset
housing = pd.read_csv('housing.csv')

# For this demonstration, consider only 'median_income' and 'median_house_value'
housing_selected = housing[['median_income', 'median_house_value']].copy()

# Random split: This splits the data randomly without preserving any specific distribution.
train_set_random, test_set_random = train_test_split(housing_selected, test_size=0.2, random_state=42)

# For stratified sampling, first create an income category.
housing_selected['income_cat'] = pd.cut(housing_selected['median_income'],
                                         bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                                         labels=[1, 2, 3, 4, 5])

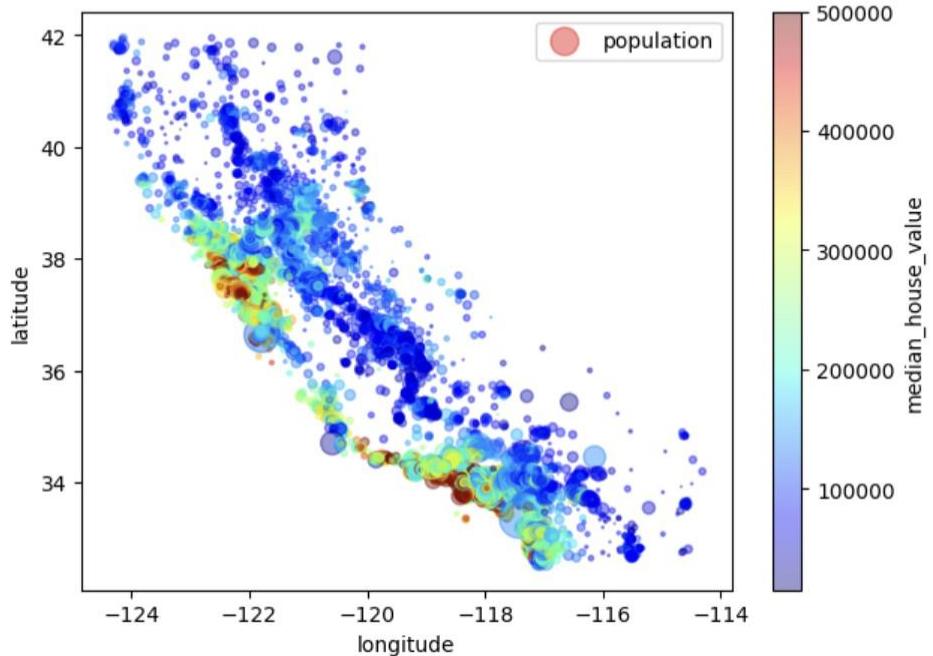
# Use StratifiedShuffleSplit to ensure the income distribution is preserved in both sets.
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing_selected, housing_selected['income_cat']):
    strat_train_set = housing_selected.loc[train_index]
    strat_test_set = housing_selected.loc[test_index]

# Remove the temporary income category attribute.
```

```
for dataset in (strat_train_set, strat_test_set):
    dataset.drop("income_cat", axis=1, inplace=True)
```

```
import matplotlib.pyplot as plt
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,
             s=housing["population"]/100, label="population", figsize=(7,5),
             c="median_house_value", cmap=plt.get_cmap("jet"), colorbar=True,
             plt.legend()
```

```
<matplotlib.legend.Legend at 0x7e55a2076b10>
```

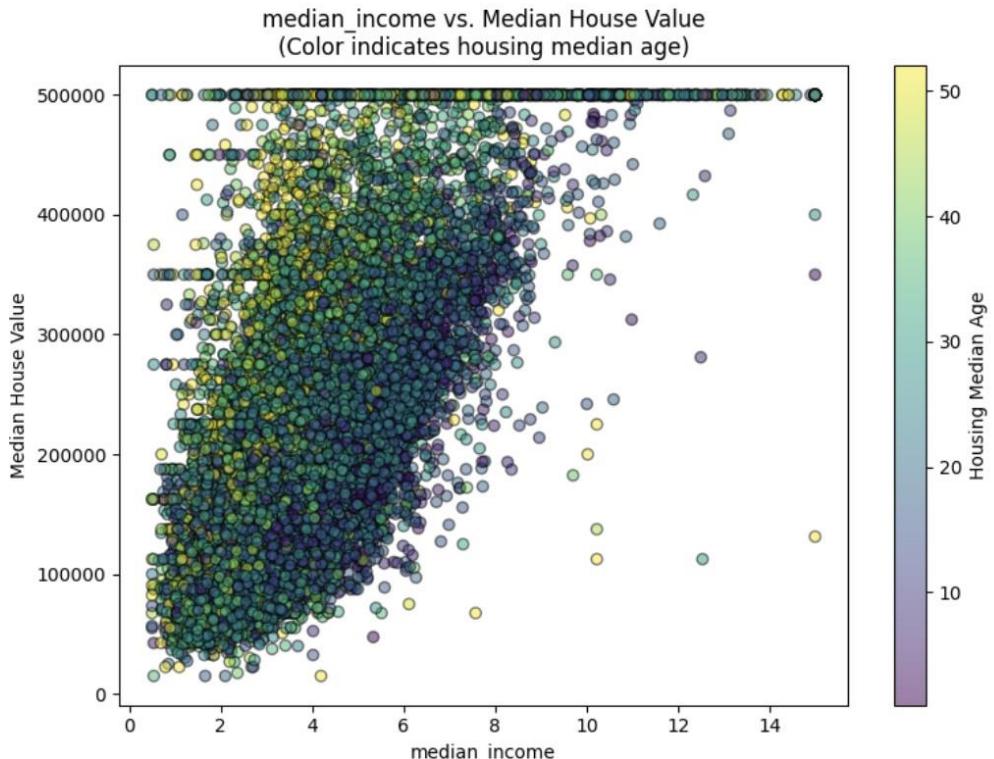


```

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,6))
# Differentiate by using 'housing_median_age' for the color
scatter = plt.scatter(housing_numeric[max_feature],
                      housing_numeric["median_house_value"],
                      alpha=0.5,
                      c=housing_numeric["housing_median_age"],
                      cmap='viridis',
                      edgecolor='k')
plt.xlabel(max_feature)
plt.ylabel("Median House Value")
plt.title(f"{max_feature} vs. Median House Value\n(Color indicates housing median age)")
# Add a colorbar to explain the color mapping
cbar = plt.colorbar(scatter)
cbar.set_label("Housing Median Age")
plt.tight_layout()
plt.show()

```



```

from sklearn.preprocessing import OneHotEncoder

# Extract the categorical attribute
housing_cat = housing[["ocean_proximity"]]

# Perform one-hot encoding
encoder = OneHotEncoder()
housing_cat_1hot = encoder.fit_transform(housing_cat).toarray()

# Create a DataFrame for the encoded features
housing_cat_1hot_df = pd.DataFrame(housing_cat_1hot,
                                    columns=encoder.get_feature_names_out(["ocean_proximity"]))
housing_cat_1hot_df.head()
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler

# Custom transformer to add engineered attributes
class CombinedAttributesAdder(BaseEstimator, TransformerMixin):

```

```

def __init__(self, add_bedrooms_per_room=True):
    self.add_bedrooms_per_room = add_bedrooms_per_room
def fit(self, X, y=None):
    return self
def transform(self, X):
    # Assumes X is a NumPy array with the following columns:
    # total_rooms (index 3), total_bedrooms (index 2), population (index 4), households (index 5)
    rooms_per_household = X[:, 3] / X[:, 5]
    population_per_household = X[:, 4] / X[:, 5]
    if self.add_bedrooms_per_room:
        bedrooms_per_room = X[:, 2] / X[:, 3]
        return np.c_[X, rooms_per_household, population_per_household, bedrooms_per_room]
    else:
        return np.c_[X, rooms_per_household, population_per_household]

# Identify numerical and categorical columns
num_attribs = housing.drop("ocean_proximity", axis=1).columns # All numeric columns
cat_attribs = ["ocean_proximity"]

# Build numerical pipeline: impute missing values, add new attributes, then scale
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', CombinedAttributesAdder()),
    ('std_scaler', StandardScaler()),
])
# Build the full pipeline combining numerical and categorical processing
full_pipeline = ColumnTransformer([
    ("num", num_pipeline, num_attribs),
    ("cat", OneHotEncoder(), cat_attribs),
])
# Process the dataset using the pipeline
housing_prepared = full_pipeline.fit_transform(housing)
print("Shape of processed data:", housing_prepared.shape)

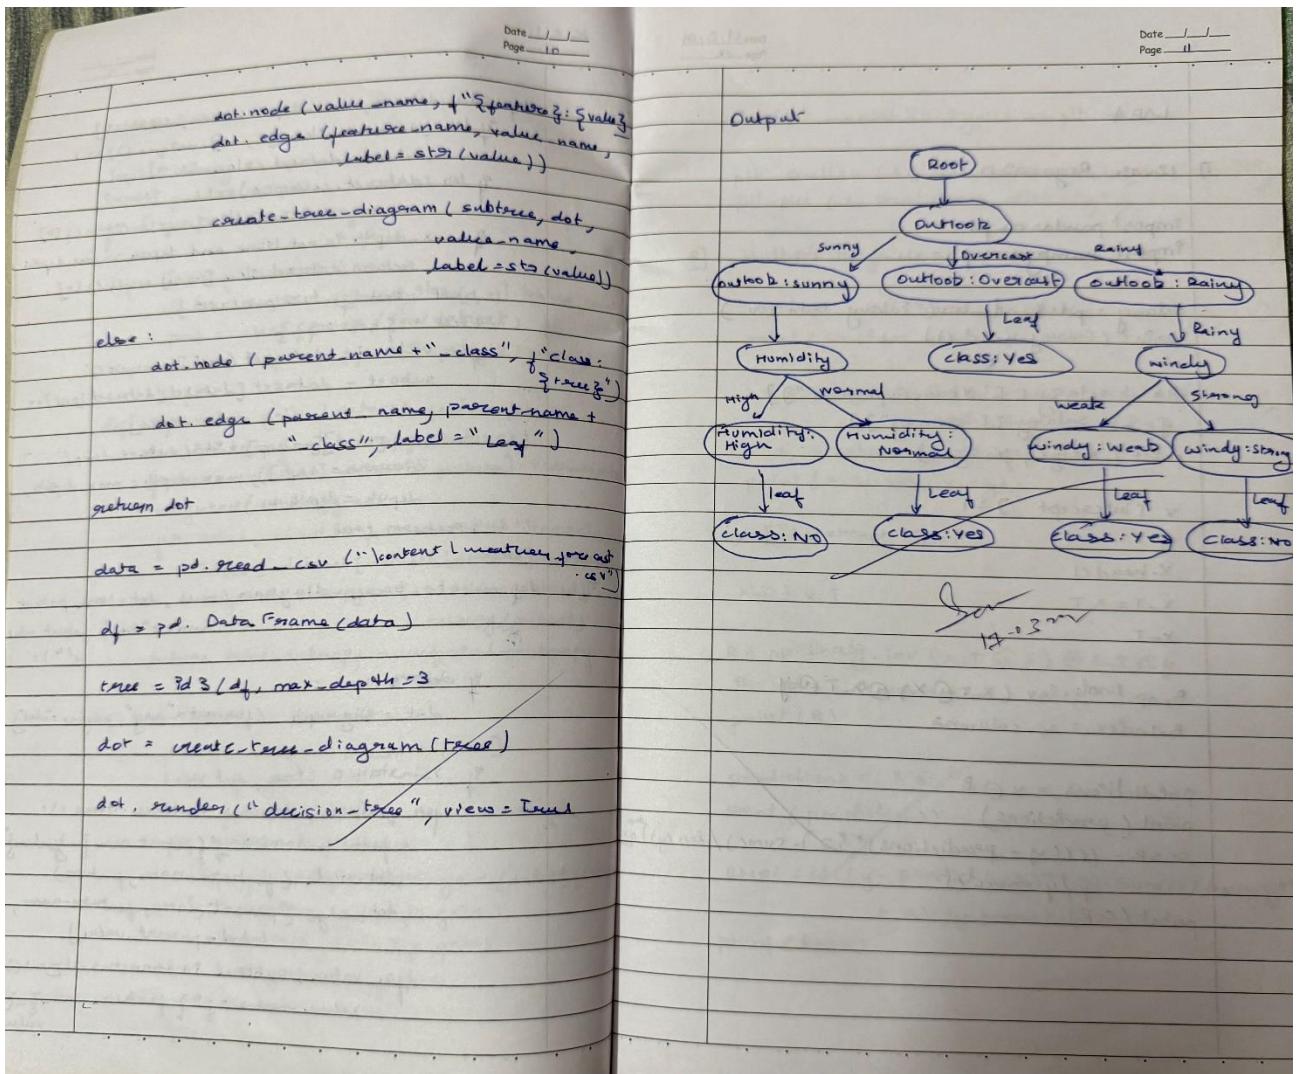
```

## LABORATORY PROGRAM – 3

**Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.**

### OBSERVATION BOOK

<p style="text-align: right;">Date 17/8/24 Page 5</p> <p><b>LAB 3</b></p> <pre> import numpy as np import pandas as pd from graphviz import Digraph  def entropy(dataset):     class_counts = dataset.iloc[:, -1].value_counts()     prob = class_counts / len(dataset)     entropy = -np.sum(prob * np.log2(prob))  def information_gain(dataset, features):     total_entropy = entropy(dataset)     feature_values = dataset[features].value_counts()     weighted_entropy = 0     for value, count in feature_values.items():         subset = dataset[dataset[features] == value]         weighted_entropy += (count / len(dataset)) * entropy(subset)     return total_entropy - weighted_entropy  def best_features(dataset):     features = dataset.columns[:-1]     best_info_gain = -1     best_feature = None     for feature in features:         info_gain = information_gain(dataset, feature)         if info_gain &gt; best_info_gain:             best_info_gain = info_gain             best_feature = feature     return best_feature. </pre>	<p style="text-align: right;">Date _____ Page 9</p> <pre> def id3(dataset, max_depth=None, depth=0):     if len(dataset.iloc[:, -1].unique()) == 1:         return dataset.iloc[:, -1]     if len(dataset.columns) == 1:         return dataset.iloc[:, -1].mode()[0]     if max_depth is not None and depth &gt; max_depth:         return dataset.iloc[:, -1].mode()[0]     best = best_features(dataset)     tree = {best: {}}     for value in dataset[best].unique():         subset = dataset[dataset[best] == value]         tree[best][value] = pd3(subset.drop(columns=[best]), max_depth=max_depth, depth=depth + 1)     return tree  def create_boxe_diagram(tree, dot=None, parent_name="Root", parent_value=""):     if dot is None:         dot = Digraph(format="png", engine="dot")     if not isinstance(tree, dict):         for feature, branches in tree.items():             feature_name = f"if {parent_name} == {feature}"             dot.node(feature_name, feature)             dot.edge(parent_name, feature_name, label=parent_value)         for value, subtree in branches.items():             value_name = f"if {feature_name} == {value}"             dot.node(value_name, value)             create_boxe_diagram(subtree, dot, parent_name=feature_name, parent_value=value)     else:         for feature, branches in tree.items():             feature_name = f"if {parent_name} == {feature}"             dot.node(feature_name, feature)             dot.edge(parent_name, feature_name, label=parent_value)             for value, subtree in branches.items():                 value_name = f"if {feature_name} == {value}"                 dot.node(value_name, value)                 create_boxe_diagram(subtree, dot, parent_name=feature_name, parent_value=value)     return dot </pre>
--	--



## CODE WITH OUTPUT

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree

# Load the iris dataset (make sure iris.csv is in the working directory)
iris = pd.read_csv("iris.csv")
# Assuming the last column is the target (species) and the rest are features.
X = iris.iloc[:, :-1]
y = iris.iloc[:, -1]

# Split data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train the Decision Tree classifier
clf_iris = DecisionTreeClassifier(criterion='entropy', random_state=42)
clf_iris.fit(X_train, y_train)

# Make predictions and evaluate the model
y_pred_iris = clf_iris.predict(X_test)
accuracy_iris = accuracy_score(y_test, y_pred_iris)
conf_matrix_iris = confusion_matrix(y_test, y_pred_iris)

print("IRIS Dataset Decision Tree Classifier")
print("Accuracy:", accuracy_iris)
print("Confusion Matrix:\n", conf_matrix_iris)
print("Classification Report:\n", classification_report(y_test, y_pred_iris))

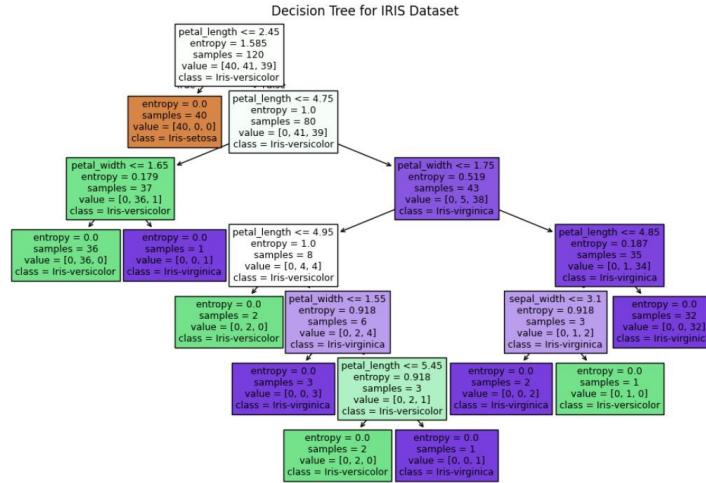
# Visualize the decision tree
plt.figure(figsize=(12, 8))
plot_tree(clf_iris, filled=True, feature_names=X.columns, class_names=clf_iris.classes_)
plt.title("Decision Tree for IRIS Dataset")
plt.show()
```

```

IRIS Dataset Decision Tree Classifier
Accuracy: 1.0
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Classification Report:
      precision    recall   f1-score   support
Iris-setosa      1.00     1.00     1.00      10
Iris-versicolor   1.00     1.00     1.00       9
Iris-virginica    1.00     1.00     1.00      11

   accuracy          1.00      30
  macro avg      1.00     1.00     1.00      30
weighted avg     1.00     1.00     1.00      30

```



```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree

# Load the drug dataset (make sure drug.csv is in the working directory)
drug = pd.read_csv("drug.csv")

# Since the target column is 'Drug', drop it from the features
X_drug = drug.drop('Drug', axis=1)
y_drug = drug['Drug']

# If there are categorical features, perform necessary encoding
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
# Encode features that are categorical
for col in X_drug.select_dtypes(include='object').columns:
    X_drug[col] = le.fit_transform(X_drug[col])
# Also encode the target variable if necessary
y_drug = le.fit_transform(y_drug)

# Split the data (80% training, 20% testing)
X_train_d, X_test_d, y_train_d, y_test_d = train_test_split(X_drug, y_drug, test_size=0.2, random_state=42)

# Initialize and train the Decision Tree classifier using entropy criterion
clf_drug = DecisionTreeClassifier(criterion='entropy', random_state=42)
clf_drug.fit(X_train_d, y_train_d)

# Make predictions and evaluate the model
y_pred_drug = clf_drug.predict(X_test_d)
accuracy_drug = accuracy_score(y_test_d, y_pred_drug)
conf_matrix_drug = confusion_matrix(y_test_d, y_pred_drug)

print("Drug Dataset Decision Tree Classifier")
print("Accuracy:", accuracy_drug)

```

```

print("Confusion Matrix:\n", conf_matrix_drug)
print("Classification Report:\n", classification_report(y_test_d, y_pred_drug))

```

```

# Visualize the decision tree
plt.figure(figsize=(12, 8))
plot_tree(clf_drug, filled=True, feature_names=X_drug.columns,
          class_names=[str(cls) for cls in clf_drug.classes_])
plt.title("Decision Tree for Drug Dataset")
plt.show()

```

```

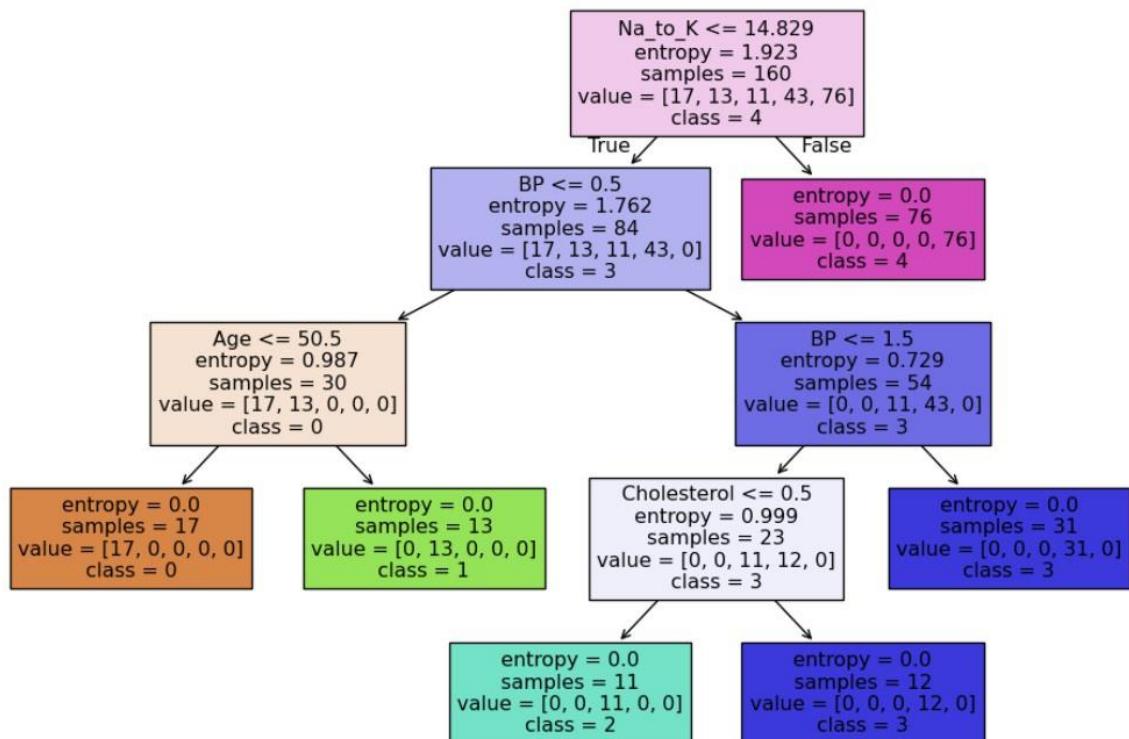
Drug Dataset Decision Tree Classifier
Accuracy: 1.0
Confusion Matrix:
[[ 6  0  0  0  0]
 [ 0  3  0  0  0]
 [ 0  0  5  0  0]
 [ 0  0  0 11  0]
 [ 0  0  0  0 15]]
Classification Report:
      precision    recall  f1-score   support

          0       1.00     1.00     1.00      6
          1       1.00     1.00     1.00      3
          2       1.00     1.00     1.00      5
          3       1.00     1.00     1.00     11
          4       1.00     1.00     1.00     15

   accuracy                           1.00      40
  macro avg       1.00     1.00     1.00      40
weighted avg       1.00     1.00     1.00      40

```

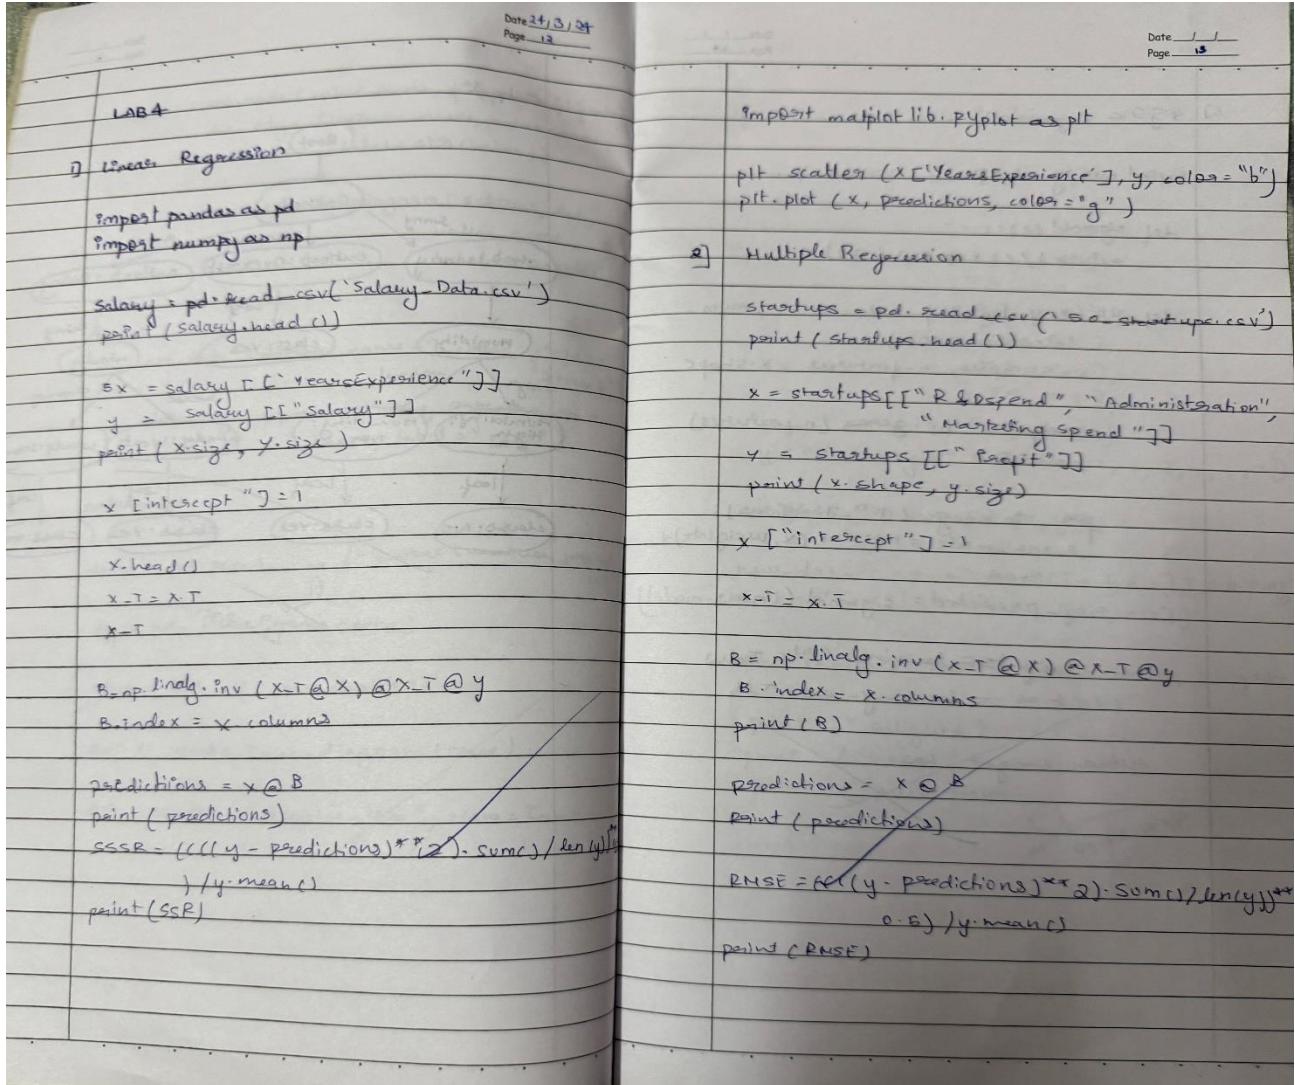
Decision Tree for Drug Dataset



## LABORATORY PROGRAM – 4

### Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

#### OBSERVATION BOOK



## CODE WITH OUTPUT

```
import pandas as pd

from sklearn.linear_model import LinearRegression
# Load the data
income_data = pd.read_csv("canada_per_capita_income.csv")
# Assumed data columns: 'Year' and 'PerCapitaIncome'
print("Canada Income Data Head:")
print(income_data.head())
# Prepare feature and target
X_income = income_data[["year"]]    # Predictor variable: Year
y_income = income_data["per capita income (US$)"]
# Build and train the linear regression model
model_income = LinearRegression()
model_income.fit(X_income, y_income)

# Predict per capita income for the year 2020
predicted_income = model_income.predict([[2020]])

print("\nPredicted per capita income for Canada in 2020:", predicted_income[0])

# Plot the data points and the regression line
plt.scatter(X_income, y_income, color='blue', label='Actual Data')
plt.plot(X_income, model_income.predict(X_income), color='red', label='Regression Line')

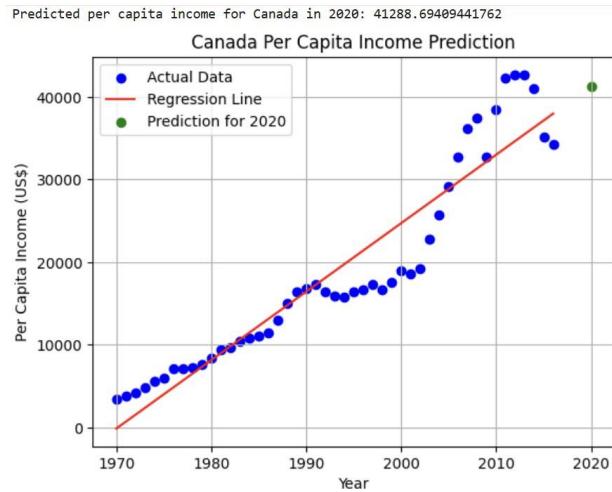
# Plot the prediction for 2020
plt.scatter(2020, predicted_income[0], color='green', label='Prediction for 2020')
```

```

# Customize the plot
plt.xlabel('Year')
plt.ylabel('Per Capita Income (US$)')
plt.title('Canada Per Capita Income Prediction')
plt.legend()
plt.grid(True)

# Display the plot
plt.show()

```



```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LinearRegression

# Load the salary data
salary_data = pd.read_csv("salary.csv")
print(salary_data.head())

# Prepare feature and target
X_salary = salary_data[["YearsExperience"]] # Predictor variable: Years of Experience
y_salary = salary_data["Salary"]

# Build and train the linear regression model
model_salary = LinearRegression()
model_salary.fit(X_salary, y_salary)

import matplotlib.pyplot as plt
# Plot the data points and the regression line
plt.scatter(X_salary, y_salary, color='blue', label='Actual Data')
plt.plot(X_salary, model_salary.predict(X_salary), color='red', label='Regression Line')

# Plot the prediction for 12 years of experience
plt.scatter(12, predicted_salary[0], color='green', label='Prediction for 12 years')

# Customize the plot
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Salary Prediction based on Experience')
plt.legend()
plt.grid(True)

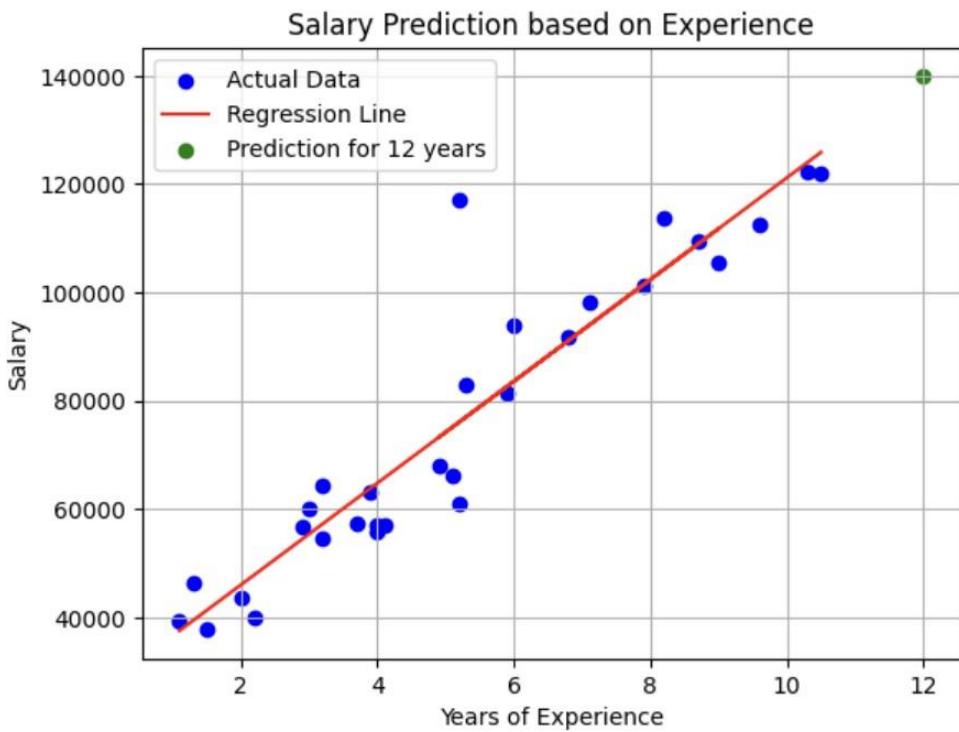
# Display the plot
plt.show()

```

---

Predicted salary for an employee with 12 years of experience: 139980.88923969213

---



```

import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression

# Read the CSV file (ensure the file is uploaded in your Colab environment)
df = pd.read_csv("hiring.csv")

# Rename columns for convenience
df.columns = ['experience', 'test_score', 'interview_score', 'salary']

print("Original Data:")
print(df)

# Function to convert experience values to numeric
def convert_experience(x):
    try:
        return float(x)
    except:
        x_lower = str(x).strip().lower()
        return num_map.get(x_lower, np.nan)

# Convert the 'experience' column using the mapping
df['experience'] = df['experience'].apply(convert_experience)

# Convert 'test_score', 'interview_score', and 'salary' to numeric (coerce errors to NaN)
df['test_score'] = pd.to_numeric(df['test_score'], errors='coerce')
df['interview_score'] = pd.to_numeric(df['interview_score'], errors='coerce')
df['salary'] = pd.to_numeric(df['salary'], errors='coerce')

print("\nData After Conversion:")
print(df)

# Fill missing values in numeric columns using the column mean
df['experience'].fillna(df['experience'].mean(), inplace=True)
df['test_score'].fillna(df['test_score'].mean(), inplace=True)
df['interview_score'].fillna(df['interview_score'].mean(), inplace=True)

print("\nData After Filling Missing Values:")
print(df)

```

```

# Prepare the feature matrix X and target vector y
X = df[['experience', 'test_score', 'interview_score']]
y = df['salary']

# Build and train the Multiple Linear Regression model
model = LinearRegression()
model.fit(X, y)

# Predict salaries for the given candidate profiles
# Candidate 1: 2 years of experience, 9 test score, 6 interview score
candidate1 = np.array([[2, 9, 6]])
predicted_salary1 = model.predict(candidate1)

# Candidate 2: 12 years of experience, 10 test score, 10 interview score
candidate2 = np.array([[12, 10, 10]])
predicted_salary2 = model.predict(candidate2)

print("\nPredicted Salary for Candidate (2 yrs, 9 test, 6 interview): $", round(predicted_salary1[0], 2))
print("Predicted Salary for Candidate (12 yrs, 10 test, 10 interview): $", round(predicted_salary2[0], 2))

import matplotlib.pyplot as plt

# Create the plot
plt.figure(figsize=(10, 6)) # Adjust figure size for better visualization
plt.scatter(df['experience'], y, color='blue', label='Actual Salary') # Plot actual salary against years of experience

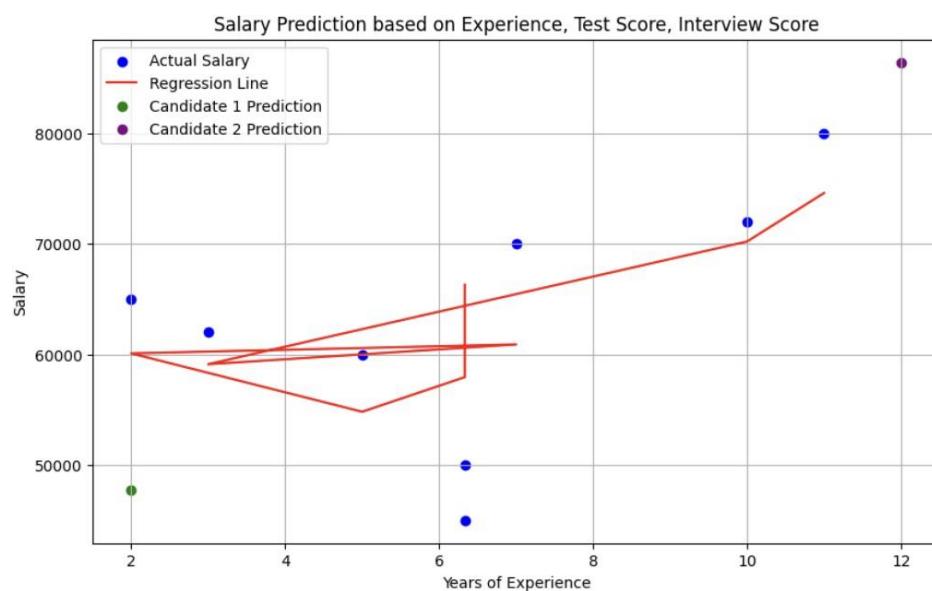
# Plot the regression line (this is an approximation since it's a multi-variable regression)
# You can visualize a single feature against the predicted salary
plt.plot(df['experience'], model.predict(X), color='red', label='Regression Line')

# Highlight predictions
plt.scatter(candidate1[0, 0], predicted_salary1, color='green', label='Candidate 1 Prediction')
plt.scatter(candidate2[0, 0], predicted_salary2, color='purple', label='Candidate 2 Prediction')

# Add labels and title
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
plt.title("Salary Prediction based on Experience, Test Score, Interview Score")

# Add a legend
plt.legend()
plt.grid(True)
plt.show()

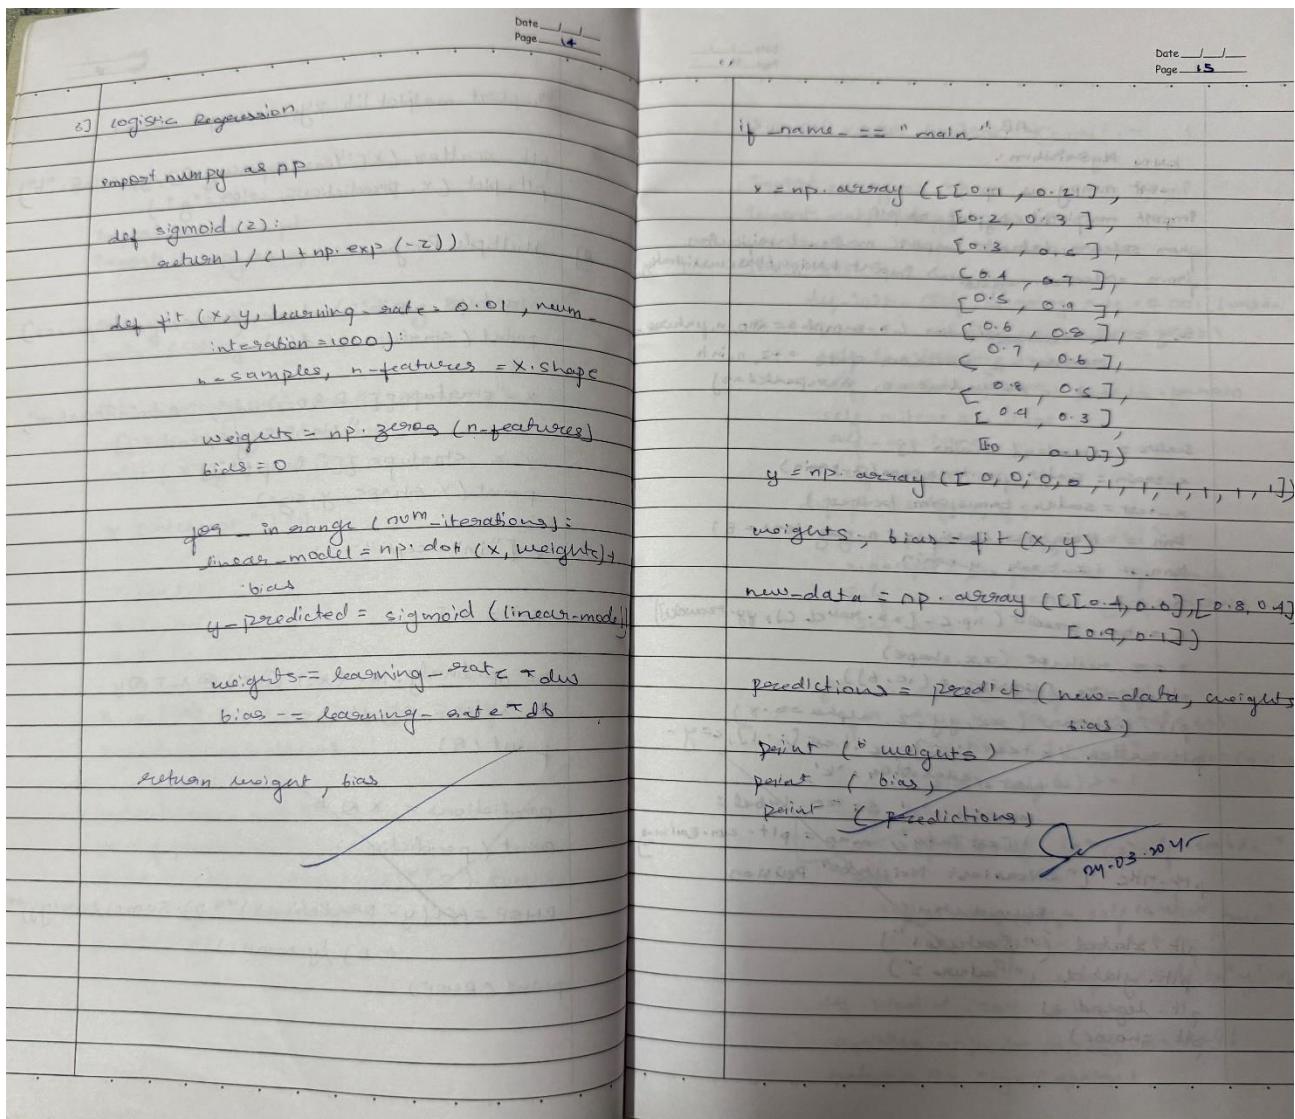
```



## LABORATORY PROGRAM – 5

### Build Logistic Regression Model for a given dataset

#### OBSERVATION BOOK



## CODE WITH OUTPUT

```
import pandas as pd
from matplotlib import pyplot as plt
# %matplotlib inline
#"%matplotlib inline" will make your plot outputs appear and be stored within the notebook.

df=pd.read_csv("insurance_data.csv")
df.head()

plt.scatter(df.age,df.bought_insurance,marker='+',color='red')

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df[['age']], df.bought_insurance, train_size=0.9, random_state=10)
X_train.shape

X_test

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()

model.fit(X_train, y_train)

X_test

y_test

y_predicted = model.predict(X_test)
y_predicted

model.score(X_test,y_test)

model.predict_proba(X_test)

y_predicted = model.predict([[60]])
y_predicted

#model.coef_ indicates value of m in y=m*x + b equation
model.coef_

#model.intercept_ indicates value of b in y=m*x + b equation
model.intercept_

#Lets defined sigmoid function now and do the math with hand
import math
def sigmoid(x):
    return 1 / (1 + math.exp(-x))

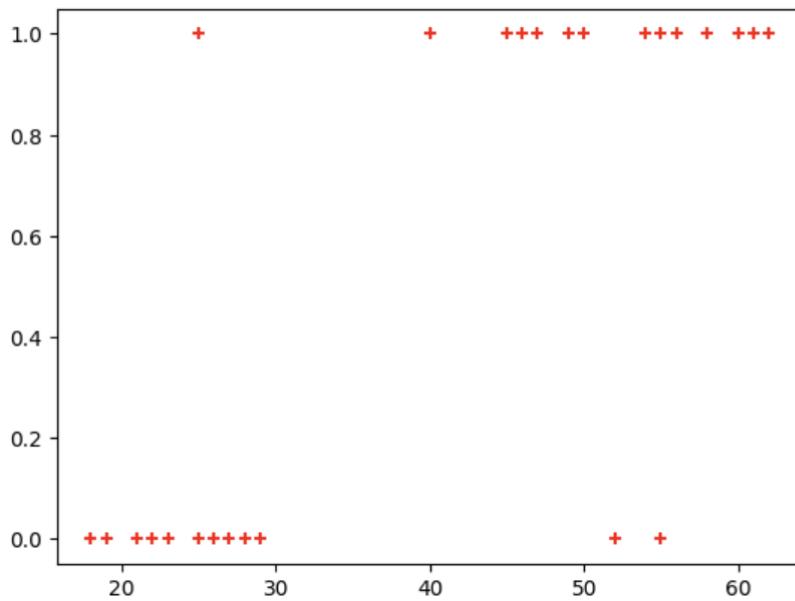
def prediction_function(age):
    z = 0.127 * age - 4.973 # 0.12740563 ~ 0.0127 and -4.97335111 ~ -4.97
    y = sigmoid(z)
    return y

age = 35
prediction_function(age)

"""0.37 is less than 0.5 which means person with 35 will not buy the insurance"""


```

'0.37 is less than 0.5 which means person with 35 will not buy the insurance'



```
# Import necessary libraries
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn import metrics
import matplotlib.pyplot as plt

# Load the Iris dataset
iris = pd.read_csv("iris.csv")
iris.head()

X=iris.drop('species',axis='columns')# Features (sepal length, sepal width, petal length, petal width)
y = iris.species # Target labels (0: Setosa, 1: Versicolor, 2: Virginica)

# Split the dataset into 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Multinomial Logistic Regression model
# Use 'multinomial' for multi-class classification and 'lbfgs' solver
model = LogisticRegression(multi_class='multinomial')

# Train the model on the training data
model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = model.predict(X_test)

# Calculate the accuracy of the model on the test data
accuracy = accuracy_score(y_test, y_pred)

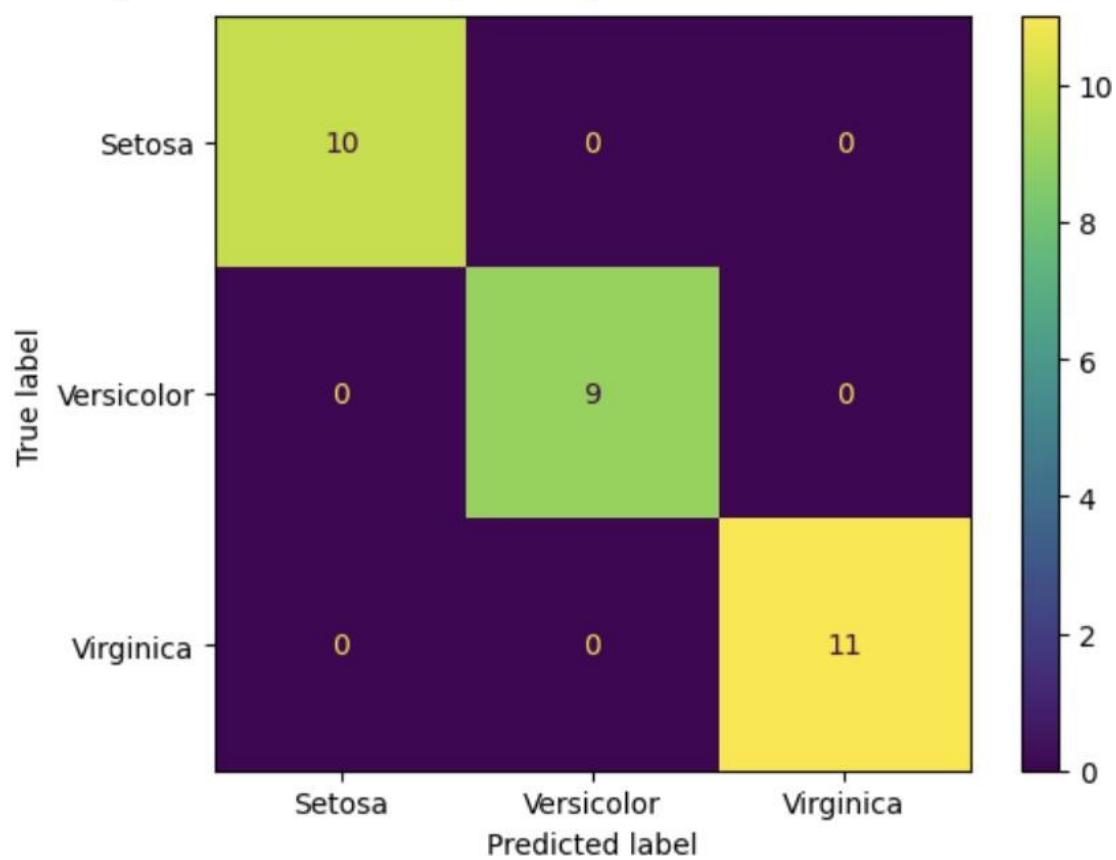
# Display the accuracy
print(f"Accuracy of the Multinomial Logistic Regression model on the test set: {accuracy:.2f}")

confusion_matrix = metrics.confusion_matrix(y_test, y_pred)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels = ["Setosa",
"Versicolor", "Virginica"])

cm_display.plot()
plt.show()
```

Accuracy of the Multinomial Logistic Regression model on the test set: 1.00



## LABORATORY PROGRAM – 6

Build KNN Classification model for a given dataset.

### OBSERVATION BOOK

<pre>def knn(self, k=Nearest Neighbors):     self.label("Feature 1")     self.ylabel("Feature 2")     self.show()  # Support Vector Machine (SVM): import numpy as np import matplotlib.pyplot as plt  class SVM:     def __init__(self, L=0.001, lambda=0.01, n_iters=1000):         self.L = L         self.b = 0         self.w = np.zeros(2)         self.n_iters = n_iters         self.w0 = None         self.b0 = None      def fit(self, X, y):         y = np.where(y &lt; 0, -1, 1)         n = len(X)         self.w = np.zeros(2)         self.b = 0         for i in range(n):             for i, j in enumerate(X):                 condition = y[i] * (np.dot(X[i], self.w))                 if condition == 0:                     self.w = self.w + self.L * (X[i] * y[i])                     self.b = self.b + y[i]                 else:                     self.w = self.w + self.L * (X[i] * y[i]) * (-1)      def predict(self, x):         affix = np.dot(x, self.w) + self.b         return np.sign(affix)</pre>	<pre>def predict(self, x):     affix = np.dot(x, self.w) + self.b     return np.sign(affix)  def visualize(self, x, y, newpoint=None):     def get_hyp(x, w, b):         return [w[0]*x[0]+w[1]*x[1]+b]/w[2]      fig = plt.figure()     for i, sample in enumerate(x):         if y[i] == 1:             plt.scatter(sample[0], sample[1], marker='o')         else:             plt.scatter(sample[0], sample[1], marker='x')      if newpoint is not None:         color = 'green' if prediction == 1 else 'orange'         plt.scatter(newpoint[0], newpoint[1], marker='x')      plt.legend()     plt.xlabel('Feature 1')     plt.ylabel('Feature 2')     plt.grid()     plt.show()</pre>
--	--

Date / /  
Page 16

LAB 5

KNN Algorithm:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.neighbors import KNeighborsClassifier

```

$x, y = \text{make_classification}(n\_samples=200, n\_features=2, n\_classes=2, n\_random_state=42, n\_in=2, n\_redundant=0, n\_separable=0)$

scaler = StandardScaler()

$x\_train = \text{scaler}.fit\_.transform(x\_train)$

$x\_test = \text{scaler}.transform(x\_test)$

knn = KNeighborsClassifier(n\_neighbors=3)

knn.fit(x\\_train, y\\_train)

$b = 0.2$

$z = \text{knn}.predict(\text{np}.c\_[x\_train], yy\_train))$

$z = z.reshape(x\_x.shape)$

plt.figure(figsize=(10, 6))

plt.contour(x\\_x, yy, z, alpha=0.8)

plt.scatter(x\\_train[:, 0], x\\_train[:, 1], c=y\\_train, edgecolor='k', marker='o', s=100, label='Train Data')

plt.scatter(x\\_test[:, 0], x\\_test[:, 1], c=z, edgecolor='k', marker='^', s=100, label='Test Data')

plt.title("KNearest Neighbors Decision Boundary")

plt.xlabel("Feature 1")

plt.ylabel("Feature 2")

plt.legend()

plt.show()

Date / /  
Page 17

SUPPORT VECTOR MACHINE (SVM) :-

```

import numpy as np
import matplotlib.pyplot as plt

```

class SVM:

```

def __init__(self, learning_rate=0.001, lambda_=0.01, n_iters=1000):
    self.ln = learning_rate
    self.lmbda = lambda_
    self.lmbda_param = lambda_.param
    self.n_iters = n_iters
    self.w = None
    self.b = None

```

def fit(self, x, y):

$y = \text{np}.where(y < 0, -1, 1)$

$n\_samples, n\_features = x.shape$

$\text{self}.w = \text{np}.zeros(n\_features)$

$\text{self}.b = 0$

for i in range(self.n\_iters):

for id\_x, x\_i in enumerate(x):

condition =  $y[i] * (\text{np}.dot(x_i, w) + b) > 1$

If condition:

$\text{self}.w = \text{self}.w + \frac{1}{n\_samples} (\text{y}[i] * x_i)$

else:

$\text{self}.w = \text{self}.w - \frac{1}{n\_samples} (\text{y}[i] * x_i)$

$\text{self}.b += \text{self}.ln^{-1} y[i]$

def predict(self, x):

$\text{applied} = \text{np}.dot(x, self.w) + self.b$

return np.sign(applied)

## CODE WITH OUTPUT

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# For model building and evaluation
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# ----- Part 1: IRIS Dataset ----- #
# Load the iris dataset (ensure iris.csv is in the same directory or provide correct path)
iris_df = pd.read_csv("iris.csv")

# Separate features and target
X_iris = iris_df.drop("species", axis=1)
y_iris = iris_df["species"]

# Split the data (80% training, 20% testing)
X_train_iris, X_test_iris, y_train_iris, y_test_iris = train_test_split(
    X_iris, y_iris, test_size=0.2, random_state=42
)

# Choose a value for k; here K=3 is used as an example.
knn_iris = KNeighborsClassifier(n_neighbors=3)

# Train the model on training data
knn_iris.fit(X_train_iris, y_train_iris)

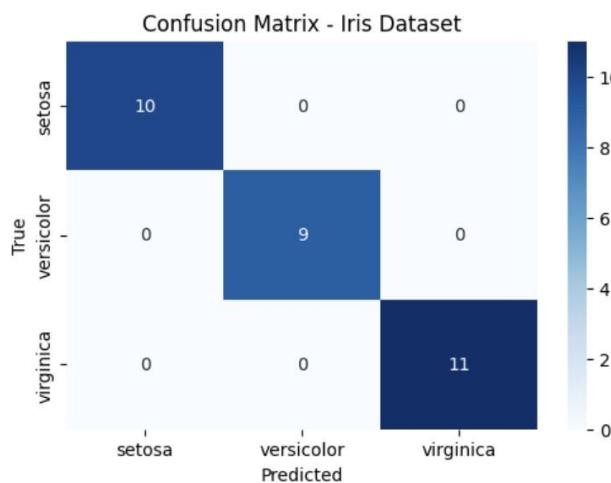
# Predict on test data
y_pred_iris = knn_iris.predict(X_test_iris)

# Calculate accuracy score
acc_iris = accuracy_score(y_test_iris, y_pred_iris)
print("IRIS Dataset Accuracy Score:", acc_iris)

# Compute confusion matrix and classification report
cm_iris = confusion_matrix(y_test_iris, y_pred_iris)
print("\nIRIS Dataset Confusion Matrix:\n", cm_iris)
```

```
cr_iris = classification_report(y_test_iris, y_pred_iris)
print("\nIRIS Dataset Classification Report:\n", cr_iris)
```

	IRIS Dataset Classification Report:			
	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30



```

# ----- Part 2: Diabetes Dataset -----#
# Load the diabetes dataset (ensure diabetes.csv is in the same directory or provide correct path)
diabetes_df = pd.read_csv("diabetes.csv")

# Separate features and target (Outcome column is assumed to be the target)
X_diabetes = diabetes_df.drop("Outcome", axis=1)
y_diabetes = diabetes_df["Outcome"]

# Perform feature scaling on the features
scaler = StandardScaler()
X_scaled_diabetes = scaler.fit_transform(X_diabetes)

# Split the scaled data (80% training, 20% testing)
X_train_diab, X_test_diab, y_train_diab, y_test_diab = train_test_split(
    X_scaled_diabetes, y_diabetes, test_size=0.2, random_state=42
)

# Choose a value for k; here K=5 is used as an example.
knn_diabetes = KNeighborsClassifier(n_neighbors=5)

# Train the model on training data
knn_diabetes.fit(X_train_diab, y_train_diab)

# Predict on test data
y_pred_diab = knn_diabetes.predict(X_test_diab)

# Calculate accuracy score
acc_diab = accuracy_score(y_test_diab, y_pred_diab)
print("Diabetes Dataset Accuracy Score:", acc_diab)

# Compute confusion matrix and classification report
cm_diab = confusion_matrix(y_test_diab, y_pred_diab)
print("\nDiabetes Dataset Confusion Matrix:\n", cm_diab)

```

```

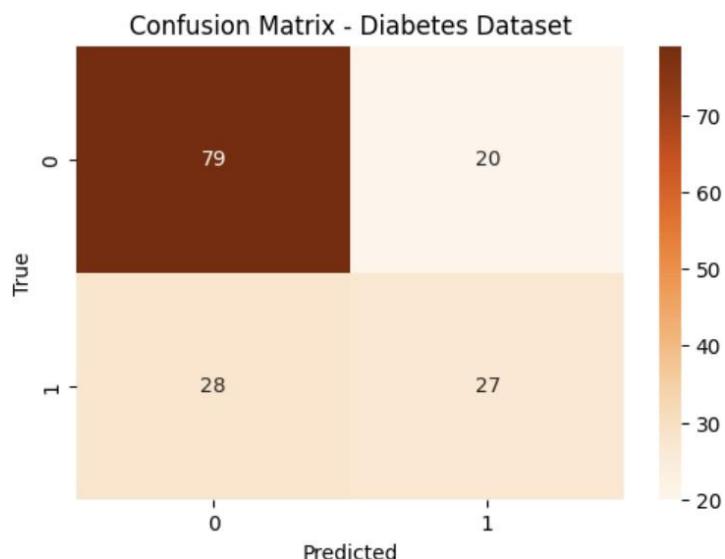
cr_diab = classification_report(y_test_diab, y_pred_diab)
print("\nDiabetes Dataset Classification Report:\n", cr_diab)

Diabetes Dataset Classification Report:
precision    recall    f1-score   support

          0       0.74      0.80      0.77      99
          1       0.57      0.49      0.53      55

   accuracy                           0.69      154
macro avg       0.66      0.64      0.65      154
weighted avg    0.68      0.69      0.68      154

```



```

# ----- Load the Dataset ----- #
# Load heart.csv (make sure the file is in your working directory)
heart_df = pd.read_csv("heart.csv")

# Display the first few rows to check the data
heart_df.head()

# ----- Data Preparation ----- #
# Separate features and target
X_heart = heart_df.drop("target", axis=1)
y_heart = heart_df["target"]

# Perform feature scaling (important for distance-based algorithms like KNN)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_heart)

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_heart, test_size=0.2, random_state=42)
# ----- Finding the Best k ----- #
# We will try a range of k values (neighbors) and select the one with maximum accuracy.
k_range = range(1, 21)
accuracy_scores = []

for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    acc = accuracy_score(y_test, y_pred)

```

```

accuracy_scores.append(acc)
print(f"k = {k} --> Accuracy: {acc:.4f}")

    k = 1 --> Accuracy: 0.8525
    k = 2 --> Accuracy: 0.8197
    k = 3 --> Accuracy: 0.8689
    k = 4 --> Accuracy: 0.8852
    k = 5 --> Accuracy: 0.9180
    k = 6 --> Accuracy: 0.9344
    k = 7 --> Accuracy: 0.9180
    k = 8 --> Accuracy: 0.8525
    k = 9 --> Accuracy: 0.8852
    k = 10 --> Accuracy: 0.8852
    k = 11 --> Accuracy: 0.8852
    k = 12 --> Accuracy: 0.8689
    k = 13 --> Accuracy: 0.8852
    k = 14 --> Accuracy: 0.8689
    k = 15 --> Accuracy: 0.9016
    k = 16 --> Accuracy: 0.8852
    k = 17 --> Accuracy: 0.8852
    k = 18 --> Accuracy: 0.9016
    k = 19 --> Accuracy: 0.8852
    k = 20 --> Accuracy: 0.8852

|: # Determine the best k value
best_k = k_range[np.argmax(accuracy_scores)]
print("\nBest k value:", best_k)

Best k value: 6

```

```

# ----- Train Final Model with Best k -----
best_knn = KNeighborsClassifier(n_neighbors=best_k)
best_knn.fit(X_train, y_train)
y_pred_best = best_knn.predict(X_test)

```

```

# Compute final accuracy, confusion matrix and classification report
final_accuracy = accuracy_score(y_test, y_pred_best)
cm = confusion_matrix(y_test, y_pred_best)
cr_text = classification_report(y_test, y_pred_best)
print("\nFinal Accuracy Score:", final_accuracy)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", cr_text)

```

Final Accuracy Score: 0.9344262295081968

Confusion Matrix:

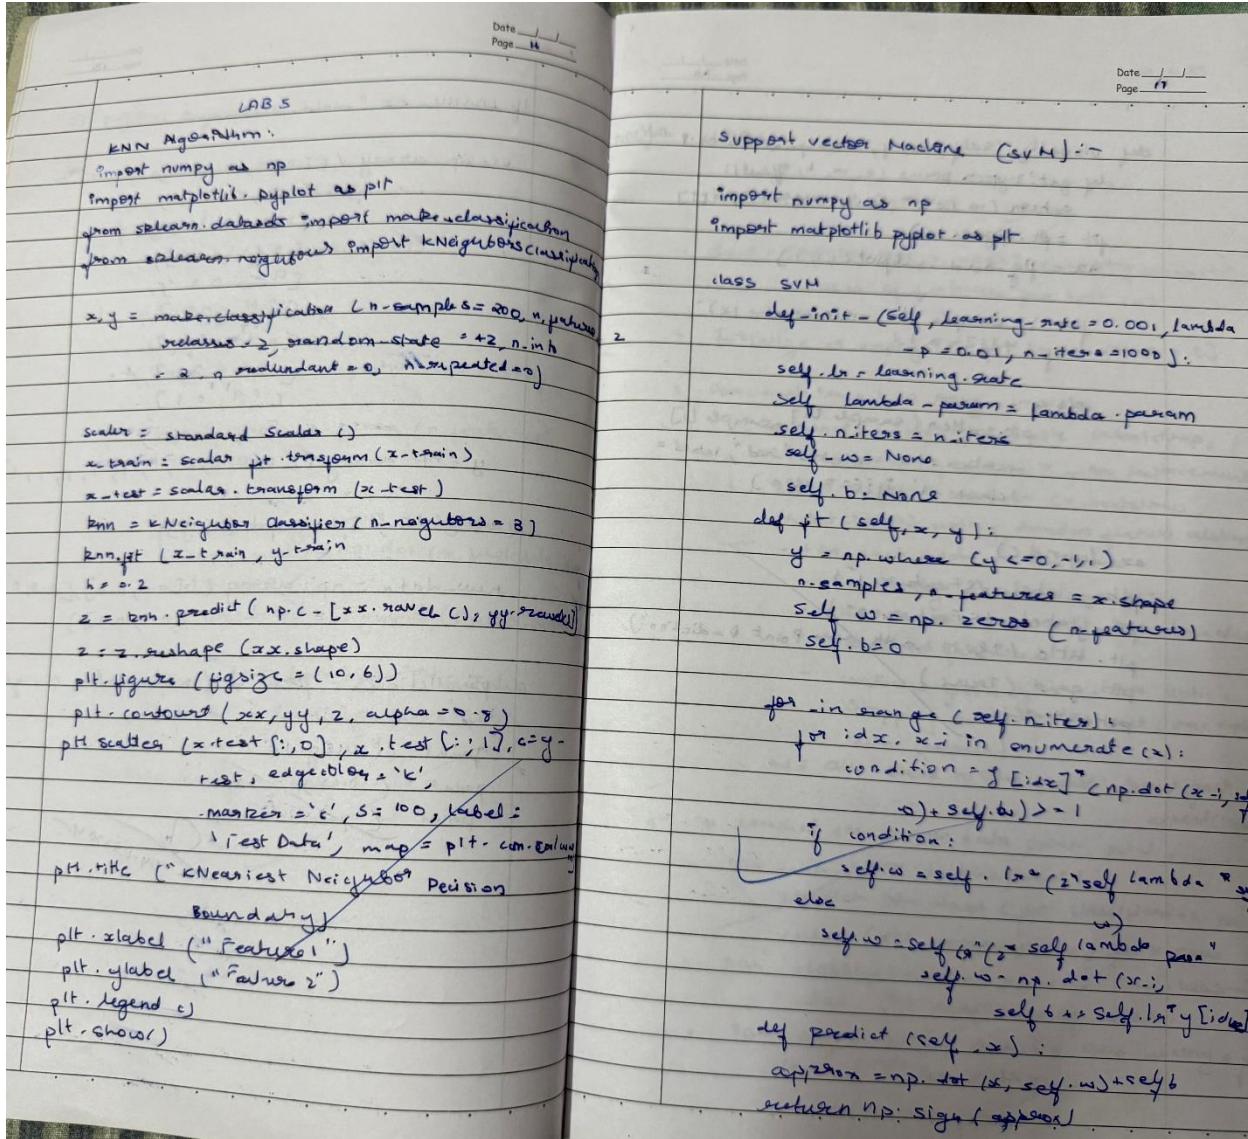
[[28 1]
[ 3 29]]

Classification Report:		precision	recall	f1-score	support
0	0.90	0.97	0.93	29	
1	0.97	0.91	0.94	32	
accuracy			0.93	61	
macro avg	0.93	0.94	0.93	61	
weighted avg	0.94	0.93	0.93	61	

# LABORATORY PROGRAM - 7

## Build Support vector machine model for a given dataset

### OBSERVATION BOOK



```
def visualize(x0, x1, y, new_point=None):
    def get_hyperplane(x, w, b, offset):
        return (w[0] * x + w[1] * x + b + offset) / w[2]
    fit = plt.figure()
    ax = fit.add_subplot(1, 1, 1)
```

```
for i, sample in enumerate(X):
    if y[i] == 1:
```

```
else:
```

```
    plt.scatter(sample[0], sample[1],
               marker='x', color="red", label=
               'class -1', if i == 0 else None)
```

```
ax.legend()
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
```

```
plt.title("SVM with New Point Prediction")
```

```
plt.grid(True)
```

```
plt.show()
```

## CODE WITH OUTPUT

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.svm import SVC

# Data points
X = np.array([[4, 1], [4, -1], [6, 0], [1, 0], [0, 1], [0, -1]])
y = np.array([1, 1, 1, -1, -1, -1])

# Fit linear SVM with a very large C to approximate hard-margin
clf = SVC(kernel='linear', C=1e6)
clf.fit(X, y)

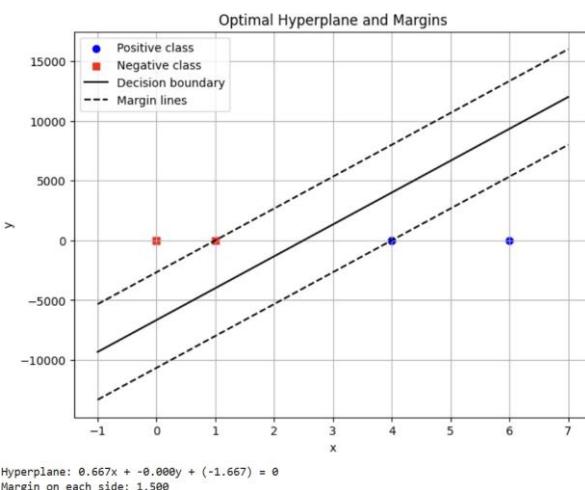
# Extract model parameters
w = clf.coef_[0]
b = clf.intercept_[0]

# Compute decision boundary and margins
xx = np.linspace(-1, 7, 500)
yy = -(w[0] * xx + b) / w[1]

# Margin offset: distance = 1/||w||
margin = 1 / np.linalg.norm(w)
yy_down = yy - np.sqrt(1 + (w[0] / w[1])**2) * margin
yy_up = yy + np.sqrt(1 + (w[0] / w[1])**2) * margin

# Plotting
plt.figure(figsize=(8, 6))
plt.scatter(X[y == 1, 0], X[y == 1, 1], c='blue', marker='o', label='Positive class')
plt.scatter(X[y == -1, 0], X[y == -1, 1], c='red', marker='s', label='Negative class')
plt.plot(xx, yy, 'k-', label='Decision boundary')
plt.plot(xx, yy_down, 'k--', label='Margin lines')
plt.plot(xx, yy_up, 'k--')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.title('Optimal Hyperplane and Margins')
plt.grid(True)
plt.show()

# Print hyperplane equation
print(f"Hyperplane: {w[0]:.3f}x + {w[1]:.3f}y + ({b:.3f}) = 0")
print(f"Margin on each side: {margin:.3f}")
```



```
import pandas as pd
```

```
# Load both datasets
```

```

iris_df = pd.read_csv("/content/iris.csv")
# 1. IRIS DATASET - SVM with RBF and Linear Kernels
X_iris = iris_df.drop("species", axis=1)
y_iris = iris_df["species"]

# Encode labels
le_iris = LabelEncoder()
y_iris_encoded = le_iris.fit_transform(y_iris)

# Split dataset
X_train_iris, X_test_iris, y_train_iris, y_test_iris = train_test_split(X_iris, y_iris_encoded, test_size=0.2, random_state=42)

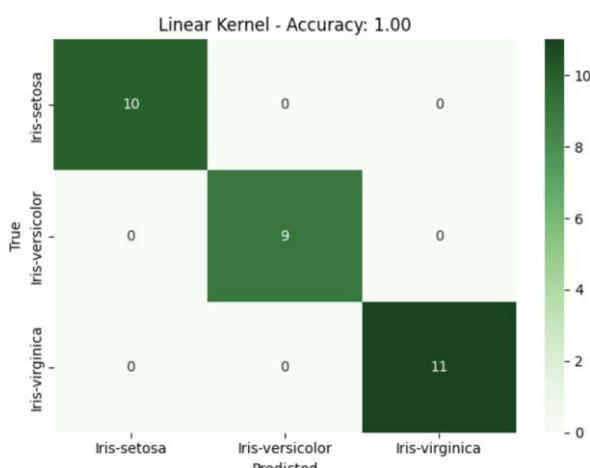
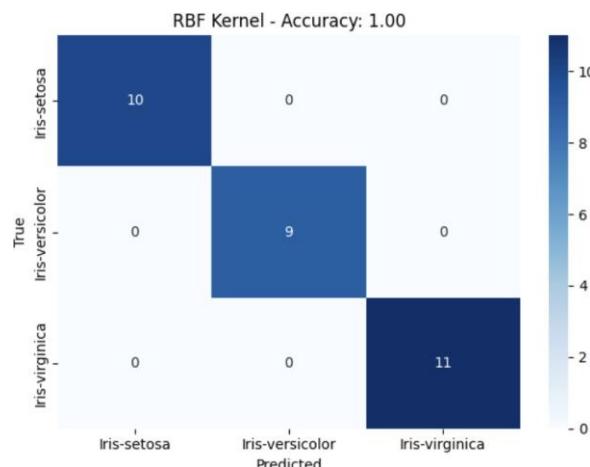
# Train models
svm_rbf = SVC(kernel='rbf')
svm_linear = SVC(kernel='linear')

svm_rbf.fit(X_train_iris, y_train_iris)
svm_linear.fit(X_train_iris, y_train_iris)

# Predictions
y_pred_rbf = svm_rbf.predict(X_test_iris)
y_pred_linear = svm_linear.predict(X_test_iris)

# Accuracy and Confusion Matrix
acc_rbf = accuracy_score(y_test_iris, y_pred_rbf)
acc_linear = accuracy_score(y_test_iris, y_pred_linear)
cm_rbf = confusion_matrix(y_test_iris, y_pred_rbf)
cm_linear = confusion_matrix(y_test_iris, y_pred_linear)

```



```

# Load dataset
letter_df = pd.read_csv("/content/letter-recognition.csv") # Update path if needed

```

```

letter_df['letter'] = LabelEncoder().fit_transform(letter_df['letter'])

# Split features and labels
X = letter_df.drop('letter', axis=1)
y = letter_df['letter']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train SVM
svm = SVC(kernel='rbf', probability=True)
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
y_prob = svm.predict_proba(X_test)

# Accuracy and Confusion Matrix
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

# ROC and AUC (one-vs-rest)
y_test_bin = label_binarize(y_test, classes=np.unique(y))
n_classes = y_test_bin.shape[1]

fpr = dict()
tpr = dict()
roc_auc = dict()

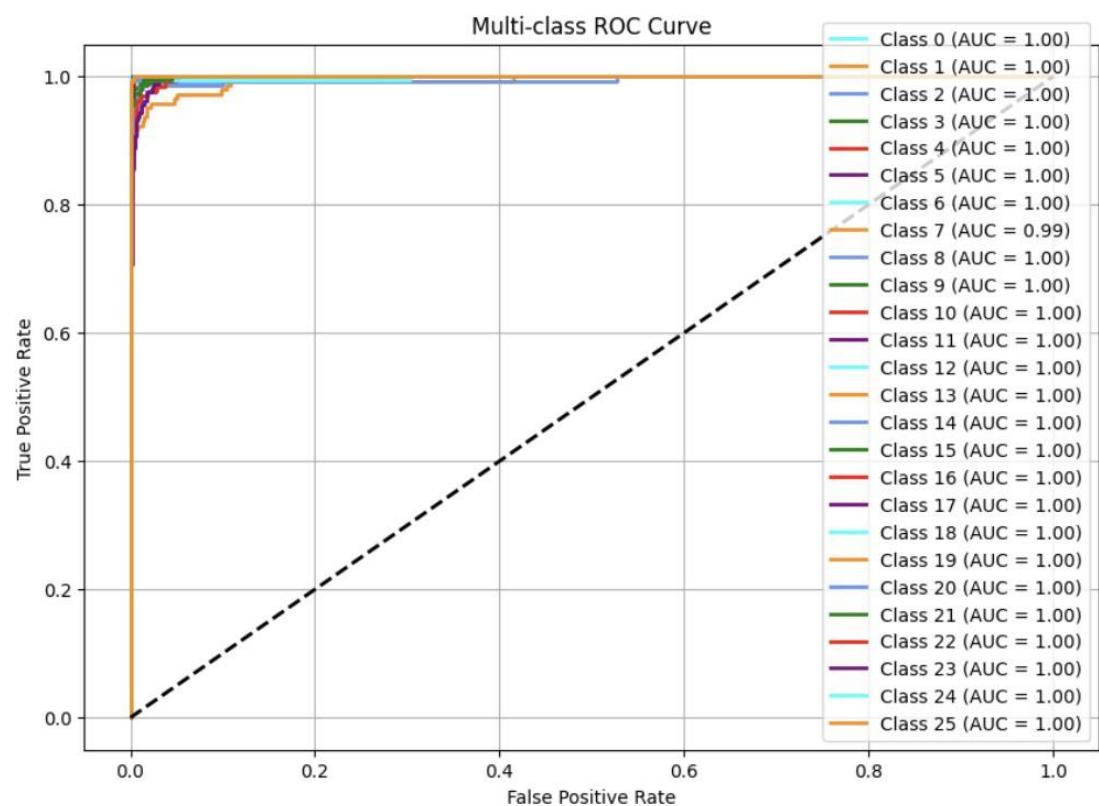
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_prob[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC Curve
plt.figure(figsize=(10, 7))
colors = cycle(['aqua', 'darkorange', 'cornflowerblue', 'green', 'red', 'purple'])

for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label=f'Class {i} (AUC = {roc_auc[i]:0.2f})')

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multi-class ROC Curve")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```



## **LABORATORY PROGRAM –8**

**Implement Random forest ensemble method on a given dataset.**

### **OBSERVATION BOOK**

Name: LAB 6

### Random Forest Algorithm

1. Input: Training data  $X$ , labels  $Y$ , number of trees ( $n$ ), number of features per split ( $k$ )
2. Initialize an empty list for trees:  $\text{forest} = []$
3. For each tree  $t$  in range ( $n$ ):
  - a. generate bootstrap sample ( $X_{\text{bootstrap}}$ ,  $Y_{\text{bootstrap}}$ ) from  $X$ ,  $Y$  with replacement
  - b. Build decision tree on ( $X_{\text{bootstrap}}$ ,  $Y_{\text{bootstrap}}$ ) using random feature selection:
    - for each node in the tree:
    - Randomly select  $k$  features
    - choose the best split based on the selected features.
    - continue growing the tree until a stopping criterion is met (max depth, min samples)
  - c. Add the tree to the forest:  $\text{forest.append(tree)}$
4. To predict for a new data point point  $x_{\text{new}}$ :
  - a. For each tree in the forest:
    - predict the class (for classification) or value (for regression) for  $x_{\text{new}}$
  - b. For classification, return the class with the majority of votes.
  - c. For regression, return the average of the tree predictions

## CODE WITH OUTPUT

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv("iris.csv") # Adjust filename if needed

# Prepare data
X = df.drop(columns=["species"]) # Assuming 'species' is the target column
y = df["species"]

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Default Random Forest with 10 trees
rf_default = RandomForestClassifier(n_estimators=10, random_state=42)
rf_default.fit(X_train, y_train)
y_pred_default = rf_default.predict(X_test)
acc_default = accuracy_score(y_test, y_pred_default)
conf_matrix_default = confusion_matrix(y_test, y_pred_default)

print(f"Default RF (10 trees) Accuracy: {acc_default}")
print("Confusion Matrix:\n", conf_matrix_default)

# Try different numbers of trees to find the best
best_acc = 0
best_n = 10
acc_list = []

for n in range(1, 101):
    rf = RandomForestClassifier(n_estimators=n, random_state=42)
    rf.fit(X_train, y_train)
    y_pred = rf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    acc_list.append((n, acc))
    if acc > best_acc:
        best_acc = acc
        best_n = n
        best_conf_matrix = confusion_matrix(y_test, y_pred)

print(f"\nBest Accuracy: {best_acc} using {best_n} trees")
print("Best Confusion Matrix:\n", best_conf_matrix)
# Plot accuracy vs number of trees
x_vals, y_vals = zip(*acc_list)
plt.plot(x_vals, y_vals, marker='o')
plt.title("Accuracy vs Number of Trees")
plt.xlabel("Number of Trees")
plt.ylabel("Accuracy")
plt.grid(True)
plt.show()
```

Default RF (10 trees) Accuracy: 1.0

Confusion Matrix:

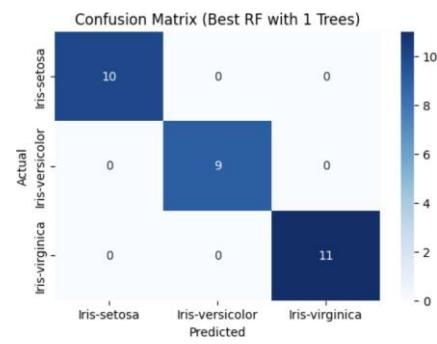
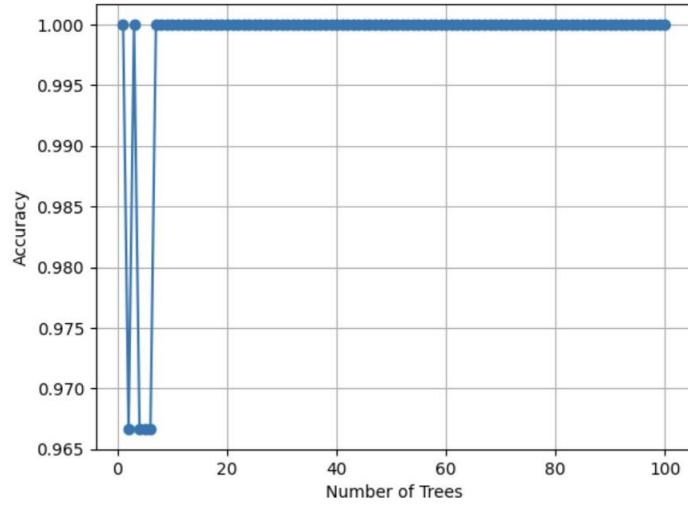
```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Best Accuracy: 1.0 using 1 trees

Best Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Accuracy vs Number of Trees



## LABORATORY PROGRAM – 9

**Implement Boosting ensemble method on a given dataset.**

### OBSERVATION BOOK

<p><i>Ada Boost Algorithm</i></p> <p>Input: Training subset <math>(x, y)</math>, number of iterations <math>t</math>  Output: Strong learners <math>H(x)</math></p> <ul style="list-style-type: none"> <li>1. Initialize sample weights: <math>w_i = 1/n</math> for each sample <math>i = 1, \dots, n</math></li> <li>2. For <math>t = 1</math> to <math>t</math>:             <ul style="list-style-type: none"> <li>a. Train weak learner <math>h_t</math> on the weighted dataset</li> <li>b. Compute error rate: <math>\epsilon_t = \sum_i w_i y_i h_t(x_i)</math>  <math>= \frac{ \{y_i \neq h_t(x_i)\} }{\sum_i w_i}</math></li> <li>c. Calculate learners weight: <math>\alpha_t = 0.5 \cdot \log(\frac{1 - \epsilon_t}{\epsilon_t})</math></li> <li>d. Update sample weights:  <math>w_i' = w_i \cdot \exp(-\alpha_t \cdot (y_i \cdot h_t(x_i)))</math> for all <math>i</math></li> <li>e. Normalize sample weights to sum to 1</li> </ul> </li> <li>3. Final prediction <math>H(x) = \text{sign}(\sum_t \alpha_t h_t(x))</math></li> <li>4. Return <math>H(x)</math></li> </ul>	<p>Date / / Page / /</p> <p><i>K-Means Algorithm</i></p> <p>Input: Data <math>X \in \mathbb{R}^{n \times d}</math> with <math>n</math> data points, number of clusters <math>k</math>  Output: cluster assignments &amp; centroids</p> <ul style="list-style-type: none"> <li>1. Randomly initialize <math>k</math> centroids <math>c_1, c_2, \dots, c_k</math></li> <li>2. Repeat until convergence:             <ul style="list-style-type: none"> <li>a. For each data point, assign it to the nearest centroid.</li> <li>- for each point <math>x_i</math>, find the closest centroid <math>c_k</math> using Euclidean distance</li> </ul> </li> <li>b. For each cluster, recompute the centroid by taking the mean of all points assigned to that cluster</li> <li>3. Return the final centroids &amp; cluster assignments.</li> </ul>
--	---

## CODE WITH OUTPUT

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay

# Load dataset
data = pd.read_csv('income.csv')

# Display basic info
print("First five rows:")
print(data.head())
print(f"\nDataset shape: {data.shape}")

# Define features and target
target_column = 'income_level'
y = data[target_column]
X = data.drop(columns=[target_column])

# Identify categorical vs numerical columns
categorical_cols = X.select_dtypes(include=['object', 'category']).columns.tolist()
numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns.tolist()
print(f"\nNumerical columns: {numerical_cols}")
print(f"Categorical columns: {categorical_cols}")

# Preprocessor: scale numericals, one-hot encode categoricals
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_cols),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
    ]
)

# Initial AdaBoost model with 10 estimators
pipeline = Pipeline([
    ('preprocess', preprocessor),
    ('clf', AdaBoostClassifier(n_estimators=10, random_state=42))
])

# Split into train/test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Train and evaluate initial model
pipeline.fit(X_train, y_train)
y_pred = pipeline.predict(X_test)
initial_acc = accuracy_score(y_test, y_pred)
print(f"Initial test accuracy (n_estimators=10): {initial_acc:.4f}")

# Hyperparameter tuning: find best n_estimators
tree_counts = list(range(10, 201, 10)) # 10,20,...,200
cv_scores = []
for n in tree_counts:
    model = Pipeline([
        ('preprocess', preprocessor),
        ('clf', AdaBoostClassifier(n_estimators=n, random_state=42))
    ])
    scores = cross_val_score(
        model, X_train, y_train, cv=5, scoring='accuracy', n_jobs=-1
    )
    mean_score = scores.mean()
```

```

cv_scores.append(mean_score)
print(f"n_estimators={n}: CV mean accuracy={mean_score:.4f}")

# Plot CV accuracy vs. number of estimators
plt.figure()
plt.plot(tree_counts, cv_scores, marker='o')
plt.title('AdaBoost CV Accuracy vs. n_estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('CV Mean Accuracy')
plt.grid(True)
plt.tight_layout()
plt.show()

# Determine optimal number of trees
best_score = max(cv_scores)
best_n = tree_counts[cv_scores.index(best_score)]
print(f"\nBest CV accuracy={best_score:.4f} with n_estimators={best_n}")

# Retrain and evaluate best model
best_model = Pipeline([
    ('preprocess', preprocess),
    ('clf', AdaBoostClassifier(n_estimators=best_n, random_state=42))
])
best_model.fit(X_train, y_train)
y_best = best_model.predict(X_test)
best_test_acc = accuracy_score(y_test, y_best)
print(f"Test accuracy with best n_estimators ({best_n}): {best_test_acc:.4f}")

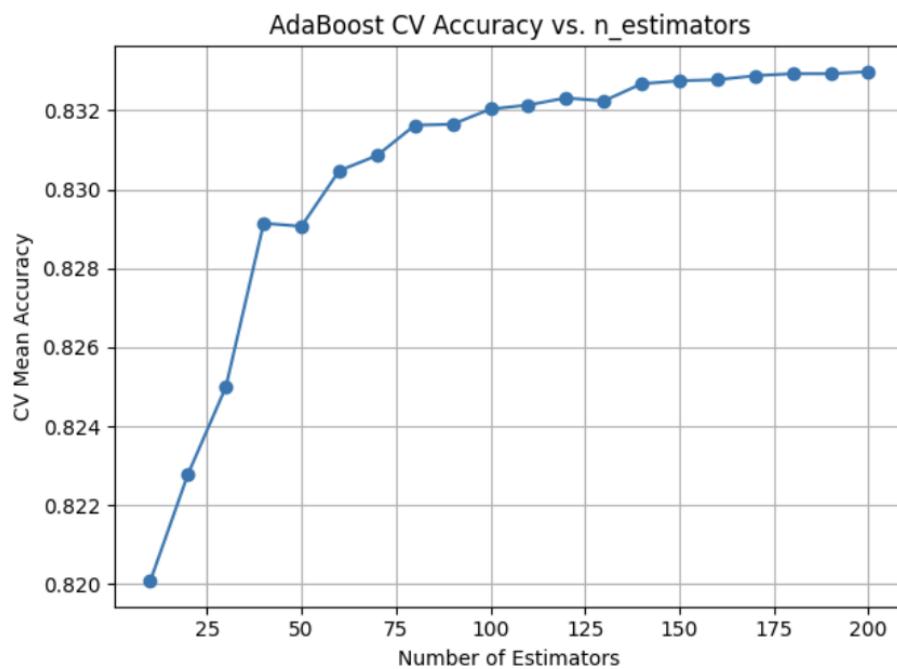
# Plot comparison of initial vs. best test accuracy
plt.figure()
plt.bar(['n=10', f'n={best_n}'], [initial_acc, best_test_acc])
plt.title('Test Accuracy: Initial vs. Optimized')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.tight_layout()
plt.show()

# Plot confusion matrix for best model
cm = confusion_matrix(y_test, y_best)
labels = best_model.named_steps['clf'].classes_
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=labels)
plt.figure()
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for Best AdaBoost Model')
plt.tight_layout()
plt.show()

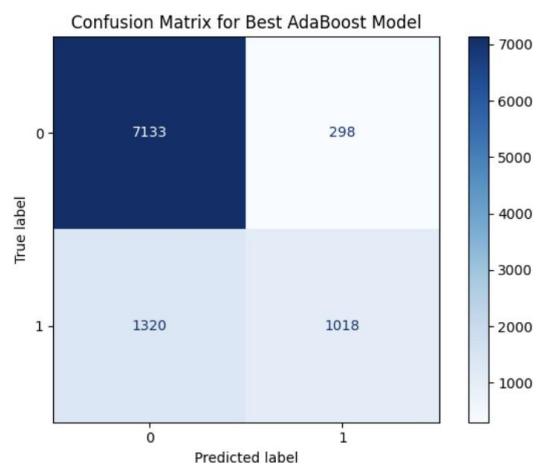
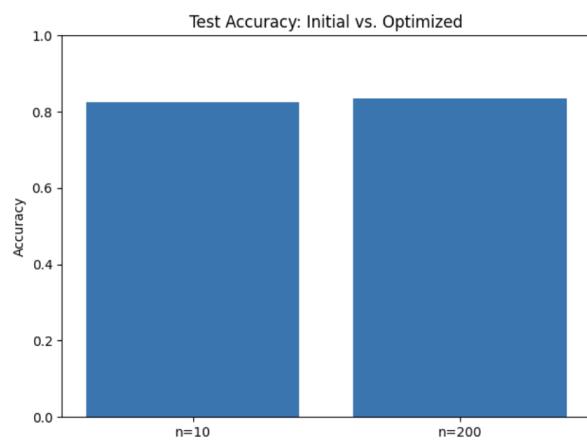
```

Dataset shape: (48842, 7)

Numerical columns: ['age', 'fnlwgt', 'education\_num', 'capital\_gain', 'capital\_loss', 'hours\_per\_week']  
Categorical columns: []  
Initial test accuracy (n\_estimators=10): 0.8257  
n\_estimators=10: CV mean accuracy=0.8201  
n\_estimators=20: CV mean accuracy=0.8228  
n\_estimators=30: CV mean accuracy=0.8250  
n\_estimators=40: CV mean accuracy=0.8291  
n\_estimators=50: CV mean accuracy=0.8291  
n\_estimators=60: CV mean accuracy=0.8305  
n\_estimators=70: CV mean accuracy=0.8309  
n\_estimators=80: CV mean accuracy=0.8316  
n\_estimators=90: CV mean accuracy=0.8316  
n\_estimators=100: CV mean accuracy=0.8320  
n\_estimators=110: CV mean accuracy=0.8321  
n\_estimators=120: CV mean accuracy=0.8323  
n\_estimators=130: CV mean accuracy=0.8322  
n\_estimators=140: CV mean accuracy=0.8327  
n\_estimators=150: CV mean accuracy=0.8327  
n\_estimators=160: CV mean accuracy=0.8328  
n\_estimators=170: CV mean accuracy=0.8329  
n\_estimators=180: CV mean accuracy=0.8329  
n\_estimators=190: CV mean accuracy=0.8329  
n\_estimators=200: CV mean accuracy=0.8330



Best CV accuracy=0.8330 with n\_estimators=200  
Test accuracy with best n\_estimators (200): 0.8344



## LABORATORY PROGRAM – 10

**Build k-Means algorithm to cluster a set of data stored in a .CSV file.**

### OBSERVATION BOOK

<b>OBSERVATION BOOK</b>	
<p><i>Ada Boost Algorithm</i></p> <p><b>Input:</b> Training subset <math>(x, y)</math>, number of iterations <math>T</math>  <b>Output:</b> Strong learners <math>H(x)</math></p> <ul style="list-style-type: none"> <li>1. Initialize sample weights: <math>w_{-i} = 1/n</math> for each sample <math>i = 1, \dots, n</math></li> <li>2. For <math>t = 1</math> to <math>T</math>:             <ul style="list-style-type: none"> <li>a. Train weak learners <math>h_t</math> on the weighted dataset</li> <li>b. Compute error rate: <math>\epsilon_t = \sum_i w_i y_i h_t(x_i)</math></li> <li>c. Calculate learners weight: <math>\alpha_t = 0.5 \log(1 - \epsilon_t) / \epsilon_t</math></li> <li>d. Update sample weights:  <math>w_{-i} = w_{-i} * \exp(-\alpha_t * (y_{-i} h_t(x_{-i}))</math> for all <math>i</math></li> <li>e. Normalize sample weights to sum to 1</li> </ul> </li> <li>3. Final prediction <math>H(x) = \text{sign}(\sum_t \alpha_t h_t(x))</math></li> <li>4. Return <math>H(x)</math></li> </ul>	<p><i>K-Mean Algorithm</i></p> <p>Date: 1/1 Page: 21</p> <p><b>Input:</b> Data <math>x \in \mathbb{R}^n</math> data points, number of clusters <math>k</math>  <b>Output:</b> cluster assignments &amp; centroids</p> <ul style="list-style-type: none"> <li>1. Randomly initialize <math>k</math> centroids <math>c_1, c_2, \dots, c_k</math></li> <li>2. Repeat until convergence:             <ul style="list-style-type: none"> <li>a. For each data point, assign it to the nearest centroid.</li> <li>- for each point <math>x_i</math>, find the closest centroid <math>c_k</math> using Euclidean distance</li> </ul> </li> <li>b. For each cluster, recompute the centroid by taking the mean of all points assigned to that cluster</li> <li>3. Return the final centroids &amp; cluster assignments.</li> </ul>

## CODE WITH OUTPUT

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

def load_data(csv_path='iris.csv'):
    """
    Try loading from csv_path; if not found, load via sklearn.
    Expects columns: sepal_length, sepal_width, petal_length, petal_width, species.
    Returns DataFrame with a 'species' column.
    """
    try:
        df = pd.read_csv(csv_path)
        # Fixed typo here: use c.strip().replace, not ace()
        df.columns = [c.strip().replace(' ', '_') for c in df.columns]
    except FileNotFoundError:
        iris = load_iris()
        df = pd.DataFrame(
            data=np.c_[iris['data'], iris['target']],
            columns=iris['feature_names'] + ['target']
        )
        df.columns = [c.strip().replace(' (cm)', "").replace(' ', '_')
                     for c in df.columns]
        df['species'] = df['target'].map(lambda x: iris['target_names'][int(x)])
    return df

def preprocess(df):
    """
    Select only petal_length & petal_width, then standard-scale.
    Returns scaled numpy array.
    """
    X = df[['petal_length', 'petal_width']].values
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)
    return X_scaled, scaler

def plot_elbow(X_scaled, max_k=10):
    """
    Compute KMeans inertia for k=1..max_k and plot the elbow curve.
    Returns list of inertias.
    """
    inertias = []
    ks = range(1, max_k + 1)
    for k in ks:
        km = KMeans(n_clusters=k, random_state=42)
        km.fit(X_scaled)
        inertias.append(km.inertia_)
    plt.figure(figsize=(6, 4))
    plt.plot(ks, inertias, 'o-', linewidth=2)
    plt.xlabel('Number of clusters (k)')
    plt.ylabel('Inertia')
    plt.title('Elbow Method for Optimal k')
    plt.xticks(ks)
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.tight_layout()
    plt.show()
    return inertias

def run_kmeans(X_scaled, k):
    """
    Fit KMeans with k clusters, return labels and fitted model.
    """
    pass
```

```

km = KMeans(n_clusters=k, random_state=42)
labels = km.fit_predict(X_scaled)
return km, labels

def plot_confusion(df, labels, k):
    """
    Builds and displays a confusion matrix comparing true species vs. cluster.
    """
    species_names = df['species'].unique()
    species_to_num = {name: idx for idx, name in enumerate(species_names)}
    true_nums = df['species'].map(species_to_num)

    cm = confusion_matrix(true_nums, labels)
    disp = ConfusionMatrixDisplay(
        confusion_matrix=cm,
        display_labels=[f"Cluster {i}" for i in range(k)])
    fig, ax = plt.subplots(figsize=(6, 6))
    disp.plot(ax=ax, cmap='Blues', colorbar=True)
    ax.set_xlabel('Predicted Cluster')
    ax.set_ylabel('True Species')
    plt.title('K-Means Clustering Confusion Matrix')
    plt.tight_layout()
    plt.show()

    cm_df = pd.DataFrame(
        cm,
        index=[f"True: {name}" for name in species_names],
        columns=[f"Cluster {i}" for i in range(k)])
    )
    print("\nConfusion Matrix (counts):")
    print(cm_df)

def main():
    # 1) Load data
    df = load_data('iris.csv')
    if 'species' not in df.columns:
        print("Error: 'species' column not found.")
        return

    # 2) Preprocess
    X_scaled, scaler = preprocess(df)

    # 3) Elbow plot to decide k
    print("Generating elbow plot to find optimal k...")
    inertias = plot_elbow(X_scaled, max_k=10)

    # 4) From the elbow you'll typically see a bend at k=3
    optimal_k = 3
    print(f"Choosing k = {optimal_k} (you can adjust this based on the plot.)")

    # 5) Run K-Means and assign clusters
    km_model, labels = run_kmeans(X_scaled, optimal_k)
    df['cluster'] = labels

    # 6) Visualize clusters in feature space
    plt.figure(figsize=(6, 4))
    plt.scatter(
        X_scaled[:, 0], X_scaled[:, 1],
        c=labels, cmap='viridis', edgecolor='k', s=50
    )
    centroids = km_model.cluster_centers_
    plt.scatter(
        centroids[:, 0], centroids[:, 1],
        marker='X', c='red', s=200, label='Centroids'
    )
    plt.xlabel('Scaled Petal Length')

```

```

plt.ylabel('Scaled Petal Width')
plt.title(f'K-Means Clusters (k={optimal_k})')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

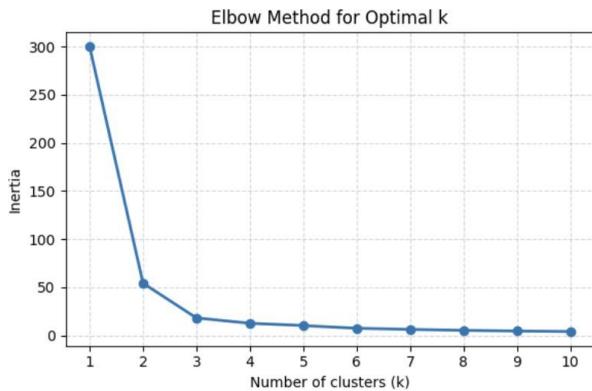
```

# 7) Confusion matrix vs. true species  
plot\_confusion(df, labels, optimal\_k)

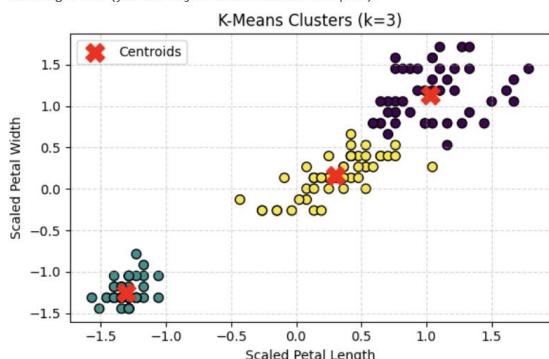
```

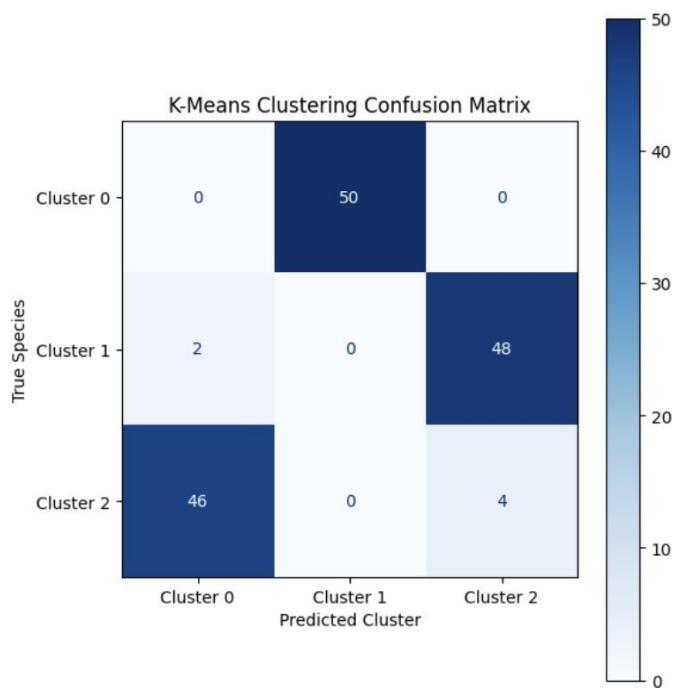
if __name__ == "__main__":
    main()

```



Choosing k = 3 (you can adjust this based on the plot).





## **LABORATORY PROGRAM – 11**

**Implement Dimensionality reduction using Principle Component Analysis  
(PCA) method.**

**OBSERVATION BOOK**

### PCA Algorithm

**Input:** Data matrix  $X$  (n samples, p features),  
number of components  $K$

**Output:** Reduced dataset  $X_{\text{PCA}}$  (n samples,  
 $K$  components)

1. Standardize the data  $X$  (subtract the mean and divide by the standard deviation)
2. calculate the covariance matrix  $C$  of the standardized data
3. compute the eigenvalues & eigenvectors of the covariance matrix
4. sort the eigenvalues in descending order  
4 choose the top  $K$  eigenvectors
5. Form the projection matrix  $P$  by selecting the corresponding eigenvectors
6. Project the data onto the new basis:  

$$X_{\text{PCA}} = X^T P$$
7. Return  $X_{\text{PCA}}$  (reduced dataset).

*Jan 19. 2024*

## CODE WITH OUTPUT

```
import pandas as pd

df = pd.read_csv("heart.csv")

# Step 3: Split Features and Target
X = df.drop("target", axis=1)
y = df["target"]

# Step 4: Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

categorical_features = ["cp", "thal", "slope"]
numerical_features = [col for col in X.columns if col not in categorical_features]

preprocessor = ColumnTransformer(transformers=[("num", StandardScaler(), numerical_features), ("cat", OneHotEncoder(), categorical_features)])
)

# Step 5: Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 6: Models
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "SVM": SVC(),
    "Random Forest": RandomForestClassifier()
}

# Step 7: Train and Evaluate Models (Before PCA)
print("Accuracy Before PCA:")
results = {}
for name, model in models.items():
    pipeline = Pipeline(steps=[("preprocessor", preprocessor), ("classifier", model)])
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f"{name}: {acc:.4f}")

from sklearn.decomposition import PCA

print("\nAccuracy After PCA (n_components=5):")
pca_results = {}

for name, model in models.items():
    pipeline_pca = Pipeline(steps=[("preprocessor", preprocessor), ("pca", PCA(n_components=5)), ("classifier", model)])
    pipeline_pca.fit(X_train, y_train)
    y_pred_pca = pipeline_pca.predict(X_test)
    acc_pca = accuracy_score(y_test, y_pred_pca)
```

```
pca_results[name] = acc_pca  
print(f"\{name}\": {acc_pca:.4f}")
```

The screenshot shows a Jupyter Notebook cell with the following output:

```
→ ┌ Accuracy Before PCA:  
    Logistic Regression: 0.9016  
    SVM: 0.8525  
    Random Forest: 0.8361  
  
└ Accuracy After PCA (n_components=5):  
    Logistic Regression: 0.8689  
    SVM: 0.8689  
    Random Forest: 0.8852
```

The code in the cell is identical to the one shown above, but the output is displayed in a more structured, collapsible format typical of Jupyter Notebooks.